

Automated Planning

EMATM0042 – Intelligent Information Systems

Thursday 14 March

kevin.mcareavey@bristol.ac.uk

Previous Lecture...

1. Agents

- Agent function
- Agent program

2. Task environments

- Problem Analysis
- Properties

3. Types of agents

- Reactive agents
- **Proactive agents**
- Cognitive agents

This Lecture...

1. Classical planning

- Planning languages

2. Planning under uncertainty

- Conformant planning
- Contingent planning
- Markov decision processes

3. Online planning

- Classical replanning

Automated Planning & Agents

- Why is automated planning important to **agent design**?
 - Generating an **agent function**?
 - Generating plans for **goal-** or **utility-based agents**?
 - Generating plans for **BDI agents**?



- Certain or uncertain **initial state**?
- (Non-)deterministic or stochastic **transitions**?
- Full, partial, or no **observability**?
- **Goals** or **utilities**?

- **Technique**?
- **Online** or **offline**?
- **Optimal** or **satisficing**?
- **Approximate**?
- **Anytime**?

- **Sequence**?
- **Partially-ordered set**?
- **Tree**?
- **Graph**?
- **Function**?

Classical Planning

Definition

- Classical planning problem $(\mathcal{S}, \mathcal{A}, s_1, T, G)$
 - Initial state $s_1 \in \mathcal{S}$
 - Deterministic transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$
 - Set of goal states $G \subseteq \mathcal{S}$
- A plan (or solution) is a finite sequence of actions (from \mathcal{A})
 - Typically interested in some notion of an optimal plan
 - A single optimal plan is sufficient




Classical Planning

Example (1)

Set of states $\mathcal{S} = \{s_1, s_2, \dots, s_8\}$

Initial state s_7

Goal states $G = \{s_1, s_2\}$

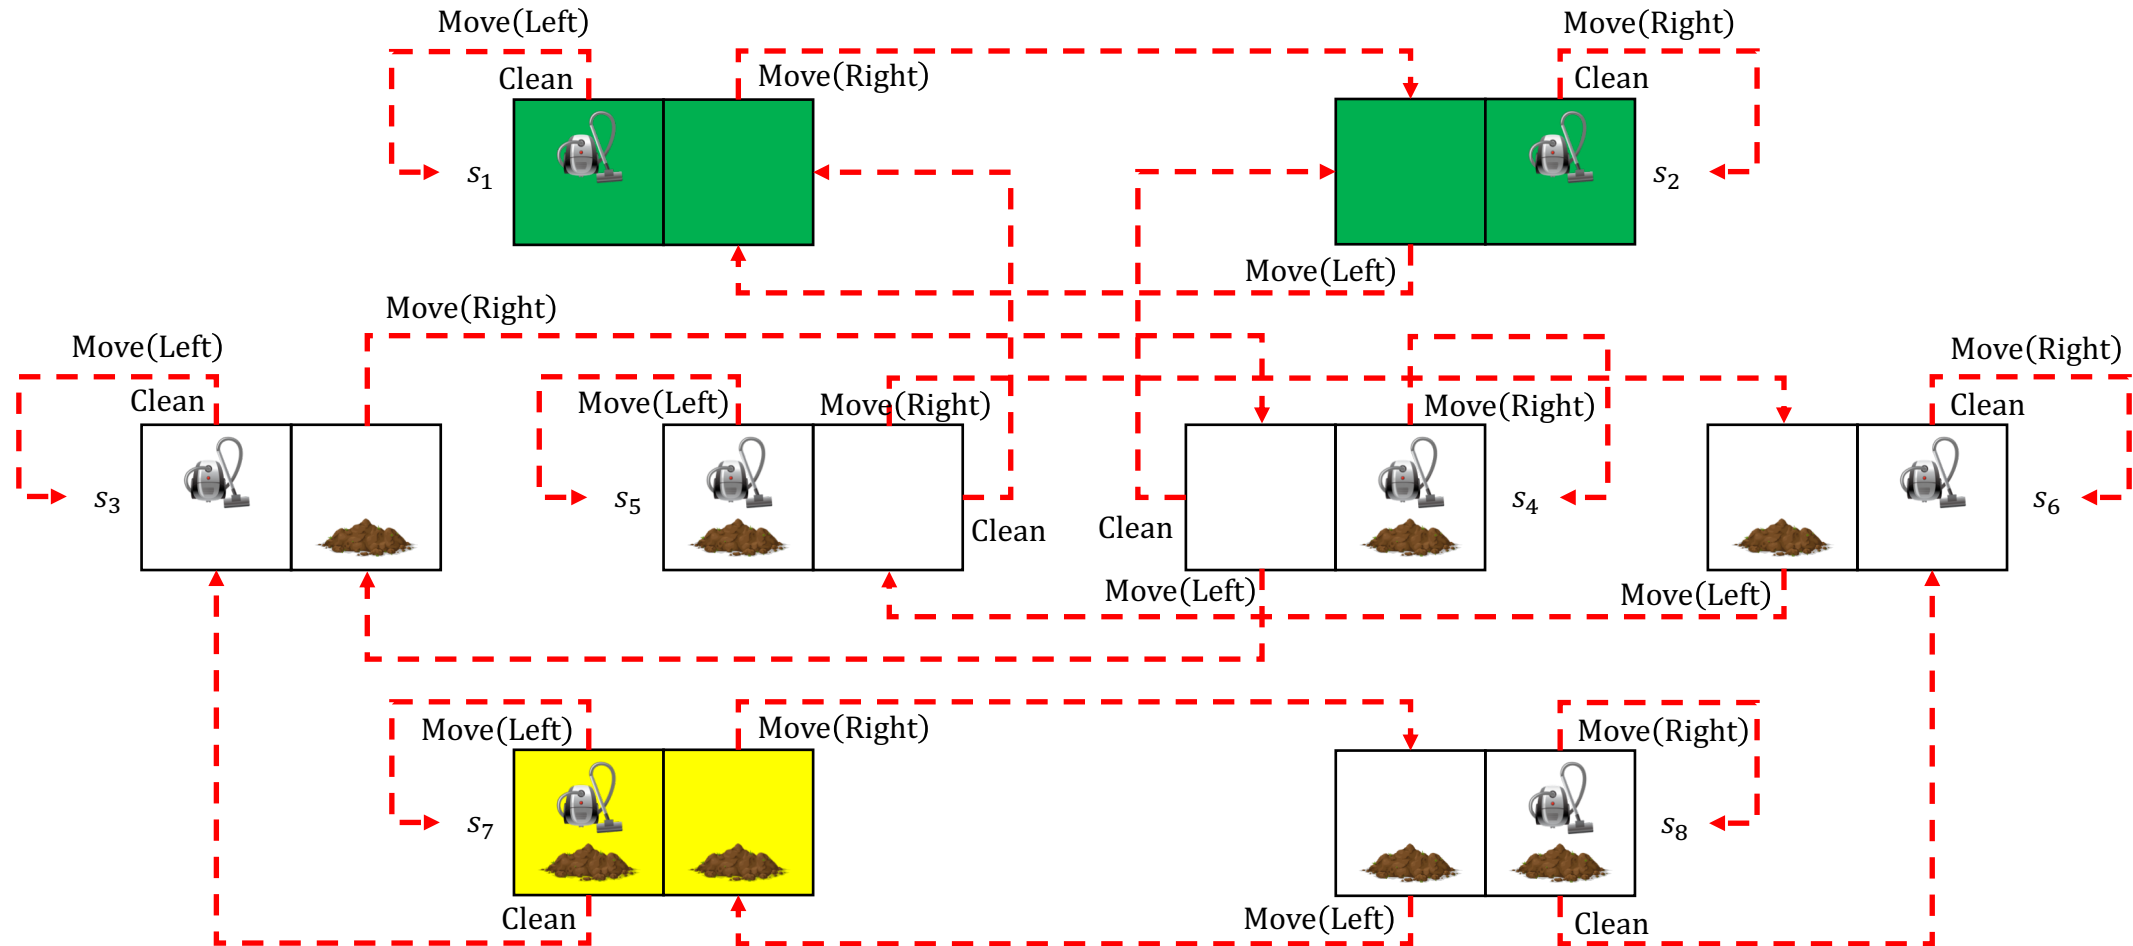


Classical Planning

Example (2)

Set of actions $\mathcal{A} = \{\text{Move(Left)}, \text{Move(Right)}, \text{Clean}\}$

Transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$

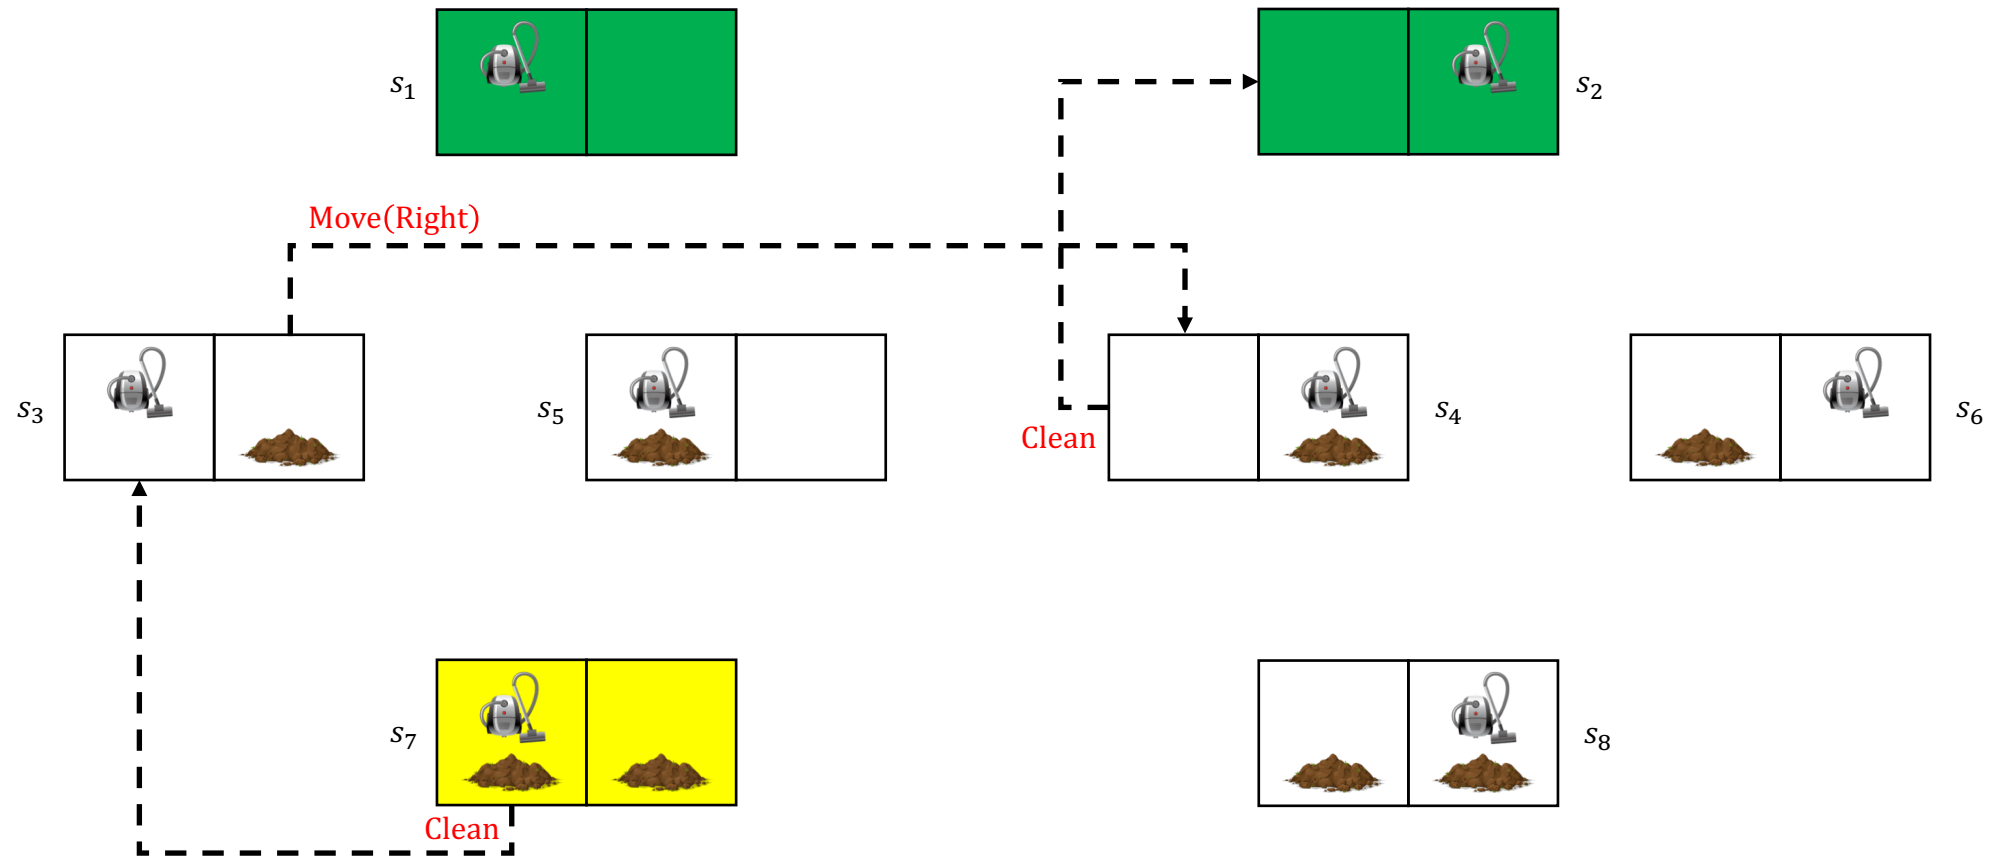


Classical Planning

Example (3)

Set of actions $\mathcal{A} = \{\text{Move(Left)}, \text{Move(Right)}, \text{Clean}\}$

Plan $\pi_1 = (\text{Clean}, \text{Move(Right)}, \text{Clean})$



Classical Planning


Solution Form (1)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$

Classical Planning

Solution Form (2)


1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



Classical Planning

Solution Form (3)


1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



Classical Planning

Solution Form (4)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$




Successor state $T(s_1, a_1) = s_2$

Classical Planning

Solution Form (5)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$




Successor state $T(s_1, a_1) = s_2$

Classical Planning

Solution Form (6)


1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



Classical Planning

Solution Form (7)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$




Successor state $T(s_2, a_2) = \dots$

Classical Planning

Solution Form (8)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$




Successor state $T(s_2, a_2) = \dots$

Classical Planning

Solution Form (9)


1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



Classical Planning

Solution Form (10)


1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



Classical Planning

Solution Form (11)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$




2. **Partial state-based policy** $\pi_{\mathcal{S}} : \mathcal{S}' \rightarrow \mathcal{A}$ with $\mathcal{S}' \subseteq \mathcal{S}$ such that $\pi_{\mathcal{S}}$ is **closed** with respect to $s_1 \in \mathcal{S}$
 - If $s \in \mathcal{S}$ is **reachable** from s_1 via $\pi_{\mathcal{S}}$ and T , then $\pi_{\mathcal{S}}$ is closed with respect s_1 if and only if $s \in \mathcal{S}'$

Classical Planning

Solution Form (12)

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



3. Policy graph $(\mathcal{S}', E, s_1, \pi_{\mathcal{S}})$

- Directed edges $E \subseteq \mathcal{S}' \times \mathcal{S}'$
- Rooted at s_1
- Nodes labelled by $\pi_{\mathcal{S}}$


2. Partial state-based policy $\pi_{\mathcal{S}} : \mathcal{S}' \rightarrow \mathcal{A}$ with $\mathcal{S}' \subseteq \mathcal{S}$ such that $\pi_{\mathcal{S}}$ is closed with respect to $s_1 \in \mathcal{S}$

- If $s \in \mathcal{S}$ is reachable from s_1 via $\pi_{\mathcal{S}}$ and T , then $\pi_{\mathcal{S}}$ is closed with respect s_1 if and only if $s \in \mathcal{S}'$

Classical Planning

Solution Form (13)

1. Finite **sequence** of actions $p = (a_1, a_2, \dots, a_n)$ from initial state $s_1 \in \mathcal{S}$



3. **Policy graph** $(\mathcal{S}', E, s_1, \pi_{\mathcal{S}})$

- Directed edges $E \subseteq \mathcal{S}' \times \mathcal{S}'$
- Rooted at s_1
- Nodes labelled by $\pi_{\mathcal{S}}$

2. Partial state-based **policy** $\pi_{\mathcal{S}} : \mathcal{S}' \rightarrow \mathcal{A}$ with $\mathcal{S}' \subseteq \mathcal{S}$ such that $\pi_{\mathcal{S}}$ is closed with respect to $s_1 \in \mathcal{S}$

- If $s \in \mathcal{S}$ is reachable from s_1 via $\pi_{\mathcal{S}}$ and T , then $\pi_{\mathcal{S}}$ is closed with respect s_1 if and only if $s \in \mathcal{S}'$

Classical Planning

Planning Languages

- STRIPS
 - States and transition functions are defined **implicitly** via **state variables** and **action descriptions**
 - **Historically important** – basis for all(?) modern planning languages
- Action Description Language (ADL)
 - Influential set of **extensions** to STRIPS language
 - Notable extensions include the introduction of **variables** (as in FOL or logic programming) and **conditional effects**
- Planning Domain Definition Language (PDDL)
 - Current **de facto standard** planning language
 - Used by the International Planning Competition (**IPC**)
 - Many **variants** (e.g. 1.0, 3.1) and **extensions** for richer planning problems (e.g. NuPDDL, PPDDL, PO-PPDDL, MA-PDDL)


Classical Planning

STRIPS

- Set of propositional **atoms** $V = \{v, v', \dots\}$ such that $\mathcal{S} = 2^V$
 - $s \models v$ where $s \in \mathcal{S}$ if and only if $v \in s$
- Set of propositional **formulas** \mathcal{L} over V with respect to $\{\neg, \wedge, \vee\}$
 - $\text{mod}(\varphi) = \{s \in \mathcal{S} \mid s \models \varphi\}$ is the **set of models of φ**
 - Language may be **restricted** (e.g. to conjunctions of positive literals)
- Initial state $s_1 \in \mathcal{S}$
 - **Set of atoms** $s_1 \subseteq V$ that are **true initially**
- Set of action descriptions $D = \{(a, \text{Pre}(a), \text{Eff}^-(a), \text{Eff}^+(a)) \mid a \in \mathcal{A}\}$
 - $\text{Pre}(a) \in \mathcal{L}$ is the **precondition** for action $a \in \mathcal{A}$
 - $\mathcal{A}(s) = \{a \in \mathcal{A} \mid s \models \text{Pre}(a)\}$ is the set of **applicable actions** in state $s \in \mathcal{S}$
 - $\text{Eff}^-(a) \subseteq V$ is the **“delete list”** for action $a \in \mathcal{A}$
 - $\text{Eff}^+(a) \subseteq V$ is the **“add list”** for action $a \in \mathcal{A}$
 - $T(s, a) = [s \setminus \text{Eff}^-(a)] \cup \text{Eff}^+(a)$ is the **successor state** after executing $a \in \mathcal{A}(s)$ in state $s \in \mathcal{S}$
 - $\text{Eff}^-(a) \cap \text{Eff}^+(a) = \emptyset$ ensures the add/delete lists are **consistent**
- Goal $\varphi \in \mathcal{L}$
 - $\text{mod}(\varphi) \subseteq \mathcal{S}$ is the set of **goal states**

Classical Planning

STRIPS – Example



s_i	Agent(Left)	Dirt(Left)	Dirt(Right)
s_1	True	False	False
s_2	False	False	False
s_3	True	False	True
s_4	False	False	True
s_5	True	True	False
s_6	False	True	False
s_7	True	True	True
s_8	False	True	True

- Initial state:
 - $\{\text{Agent(Left), Dirt(Left), Dirt(Right)}\}$
- Goal:
 - $\neg \text{Dirt(Left)} \wedge \neg \text{Dirt(Right)}$
- Action $a_1 = \text{Move(Left)}$:
 - $\text{Pre}(a_1) = \top$
 - $\text{Eff}^-(a_1) = \emptyset$
 - $\text{Eff}^+(a_1) = \{\text{Agent(Left)}\}$
- Action $a_2 = \text{Move(Right)}$:
 - $\text{Pre}(a_2) = \top$
 - $\text{Eff}^-(a_2) = \{\text{Agent(Left)}\}$
 - $\text{Eff}^+(a_2) = \emptyset$
- Action $a_3 = \text{Clean(Left)}$:
 - $\text{Pre}(a_3) = \text{Agent(Left)}$
 - $\text{Eff}^-(a_3) = \{\text{Dirt(Left)}\}$
 - $\text{Eff}^+(a_3) = \emptyset$
- Action $a_4 = \text{Clean(Right)}$:
 - $\text{Pre}(a_4) = \neg \text{Agent(Left)}$
 - $\text{Eff}^-(a_4) = \{\text{Dirt(Right)}\}$
 - $\text{Eff}^+(a_4) = \emptyset$

Classical Planning

Technique: A* State-Space Search (1)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$




Classical Planning

Technique: A* State-Space Search (2)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$




Classical Planning

Technique: A* State-Space Search (3)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



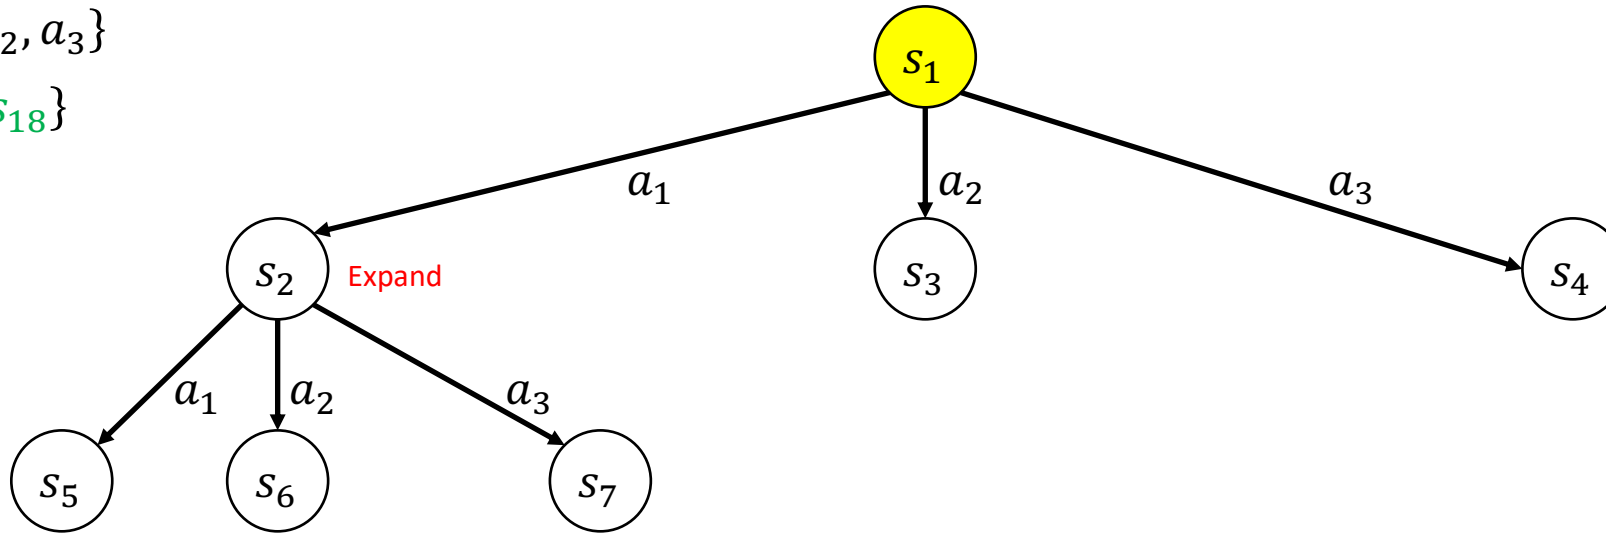
Classical Planning

Technique: A* State-Space Search (4)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



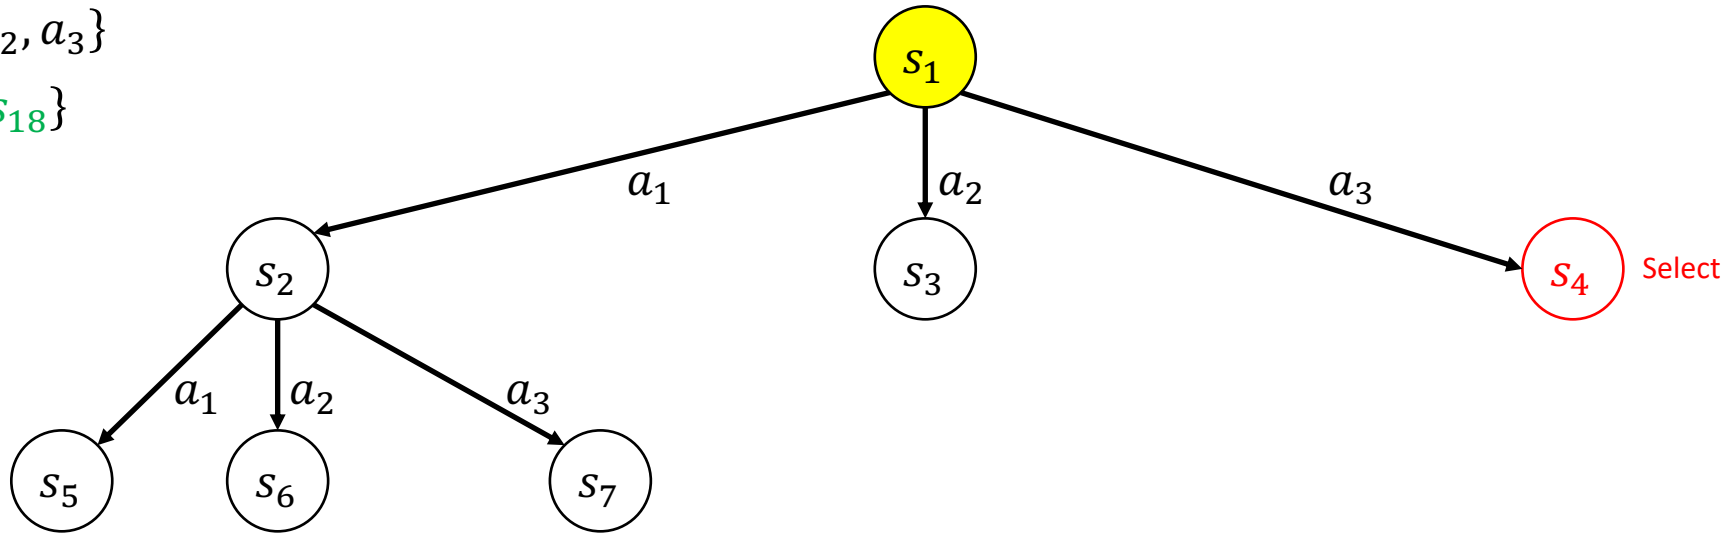
Classical Planning

Technique: A* State-Space Search (5)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



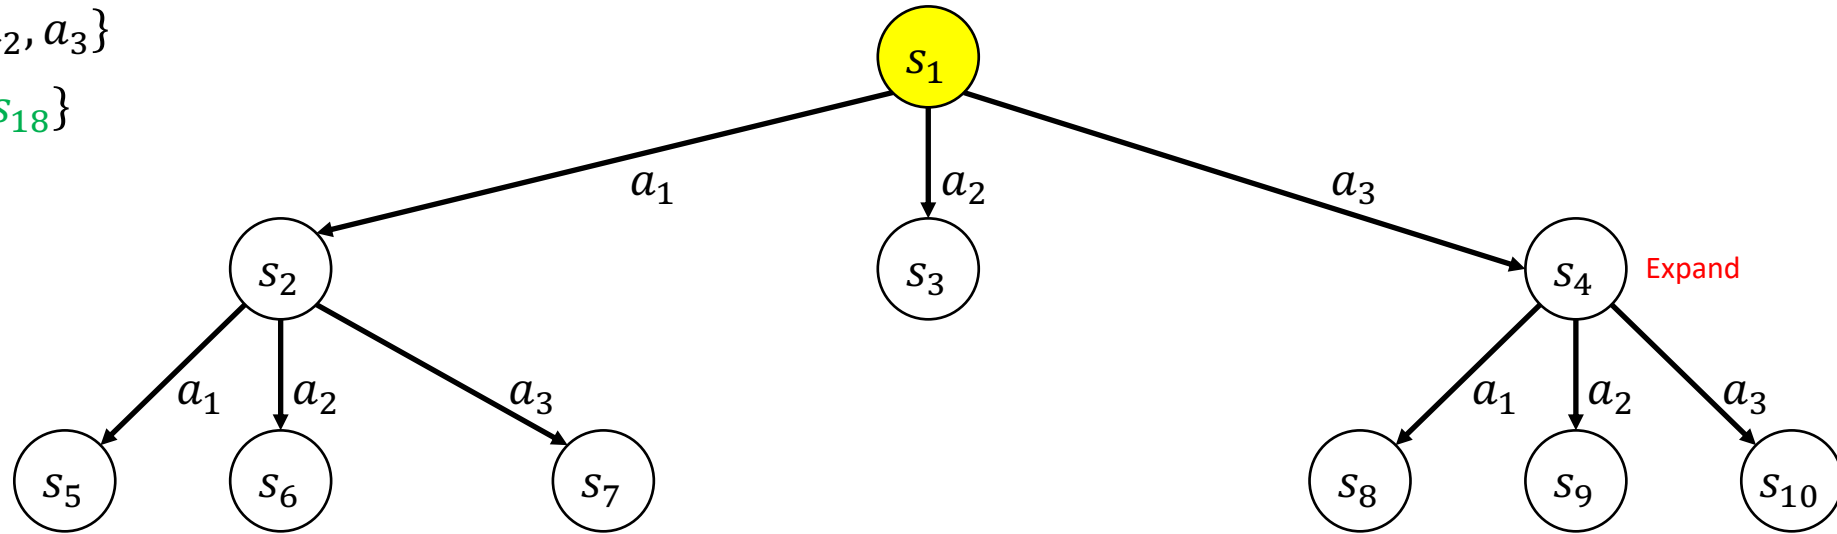
Classical Planning

Technique: A* State-Space Search (6)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



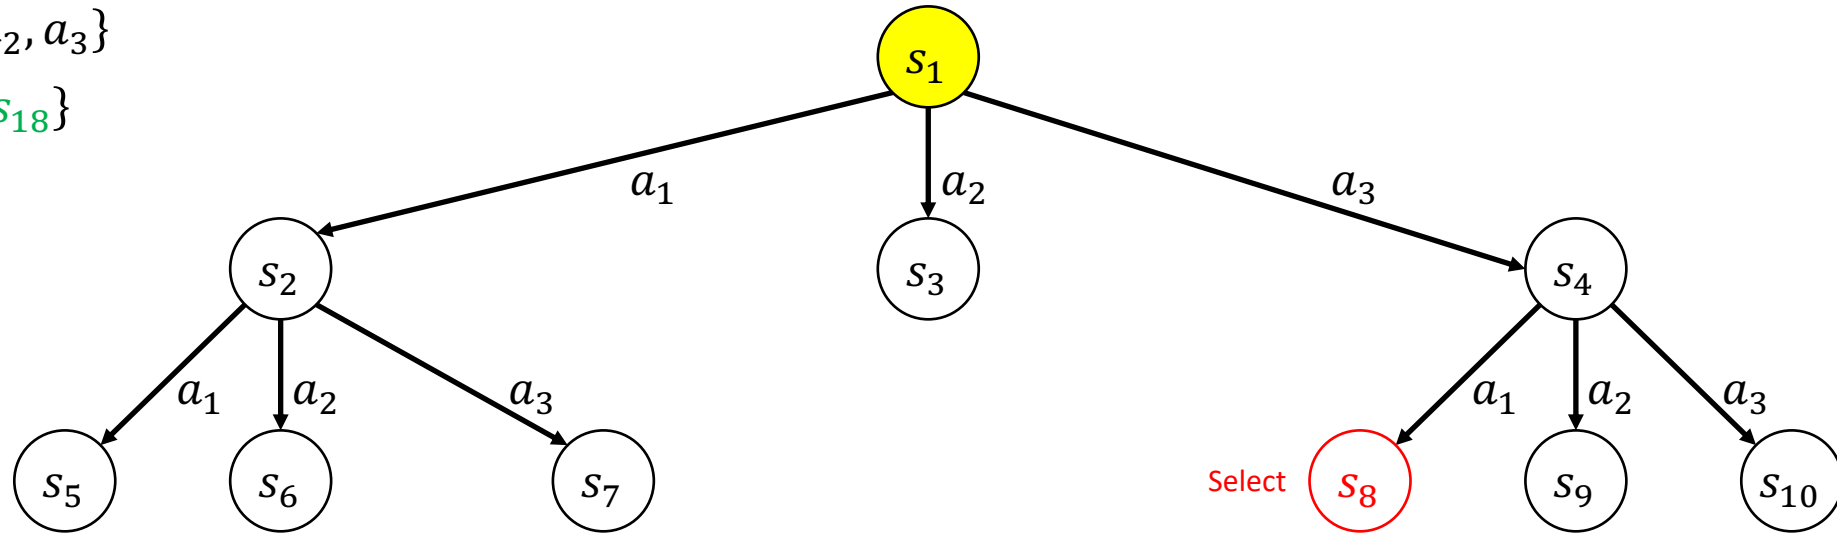
Classical Planning

Technique: A* State-Space Search (7)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



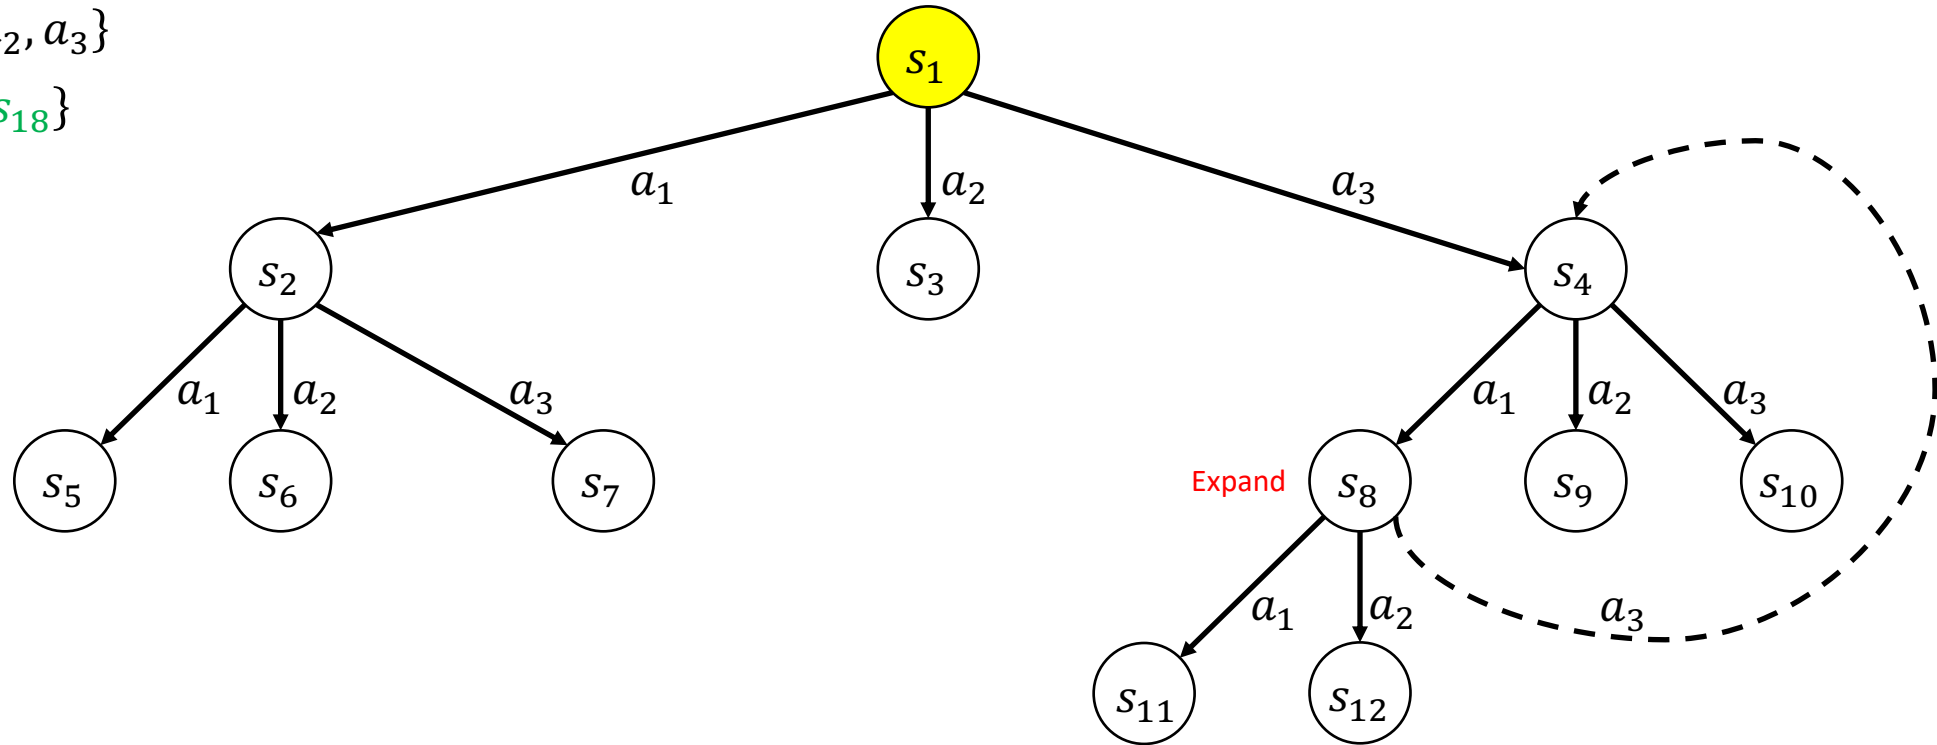
Classical Planning

Technique: A* State-Space Search (8)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



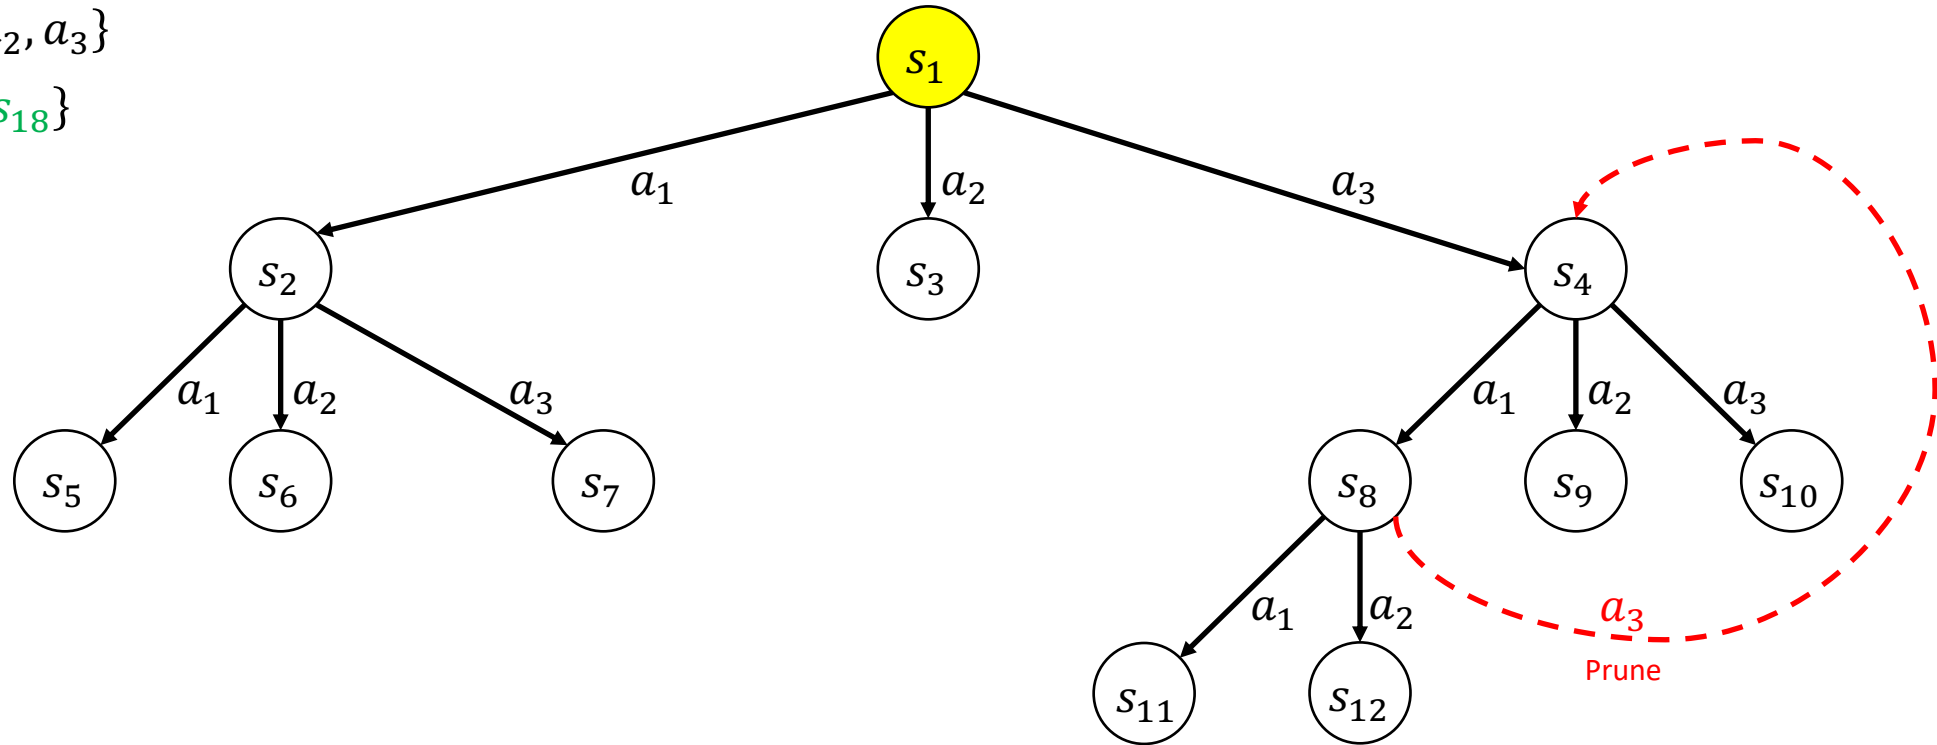
Classical Planning

Technique: A* State-Space Search (9)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



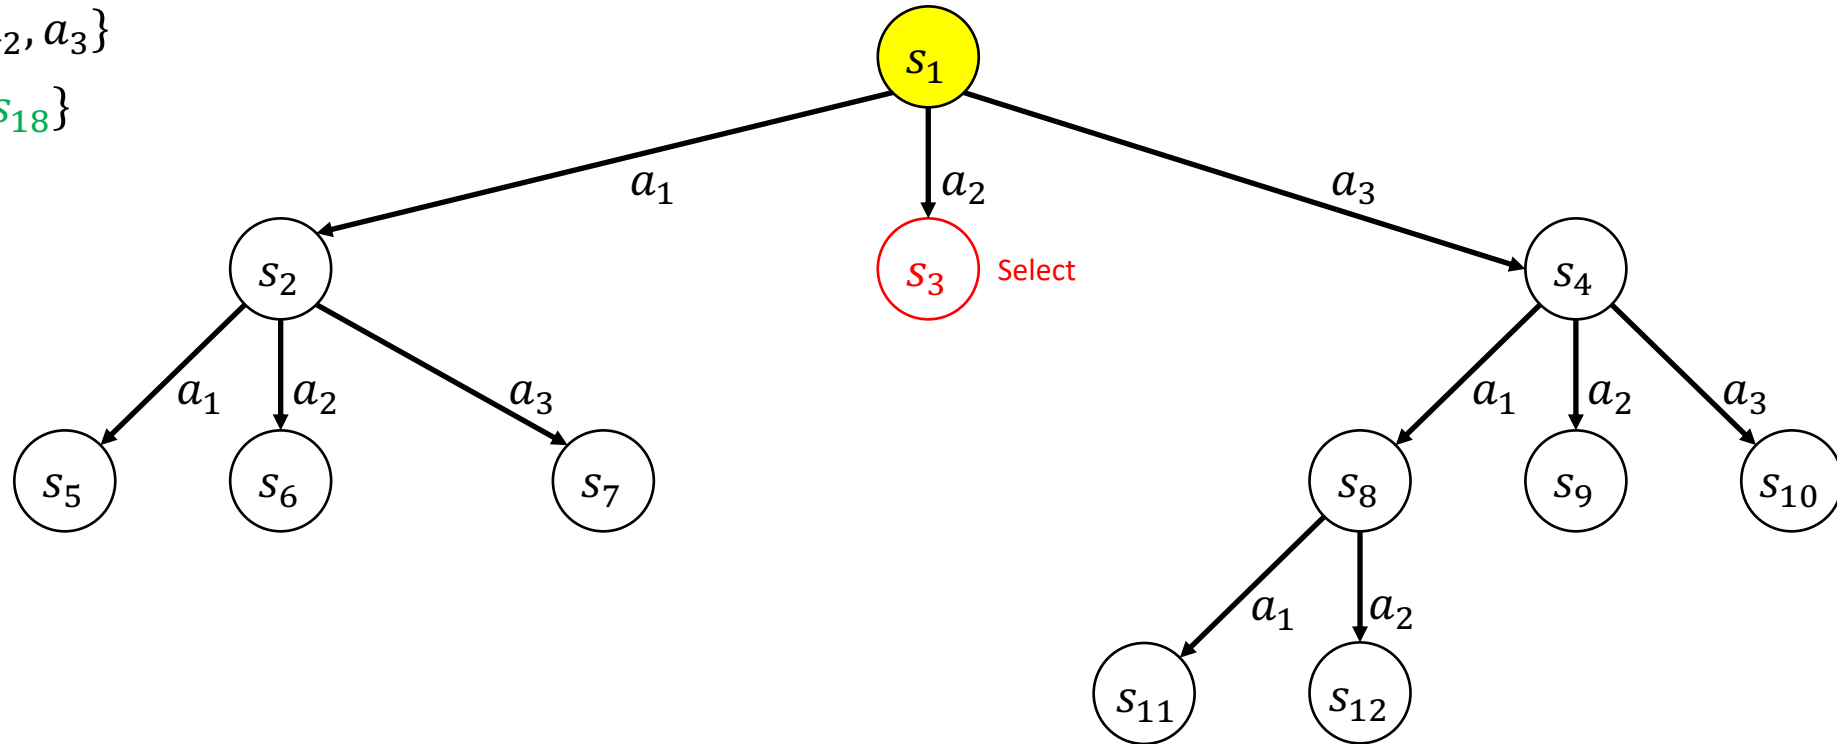
Classical Planning

Technique: A* State-Space Search (10)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



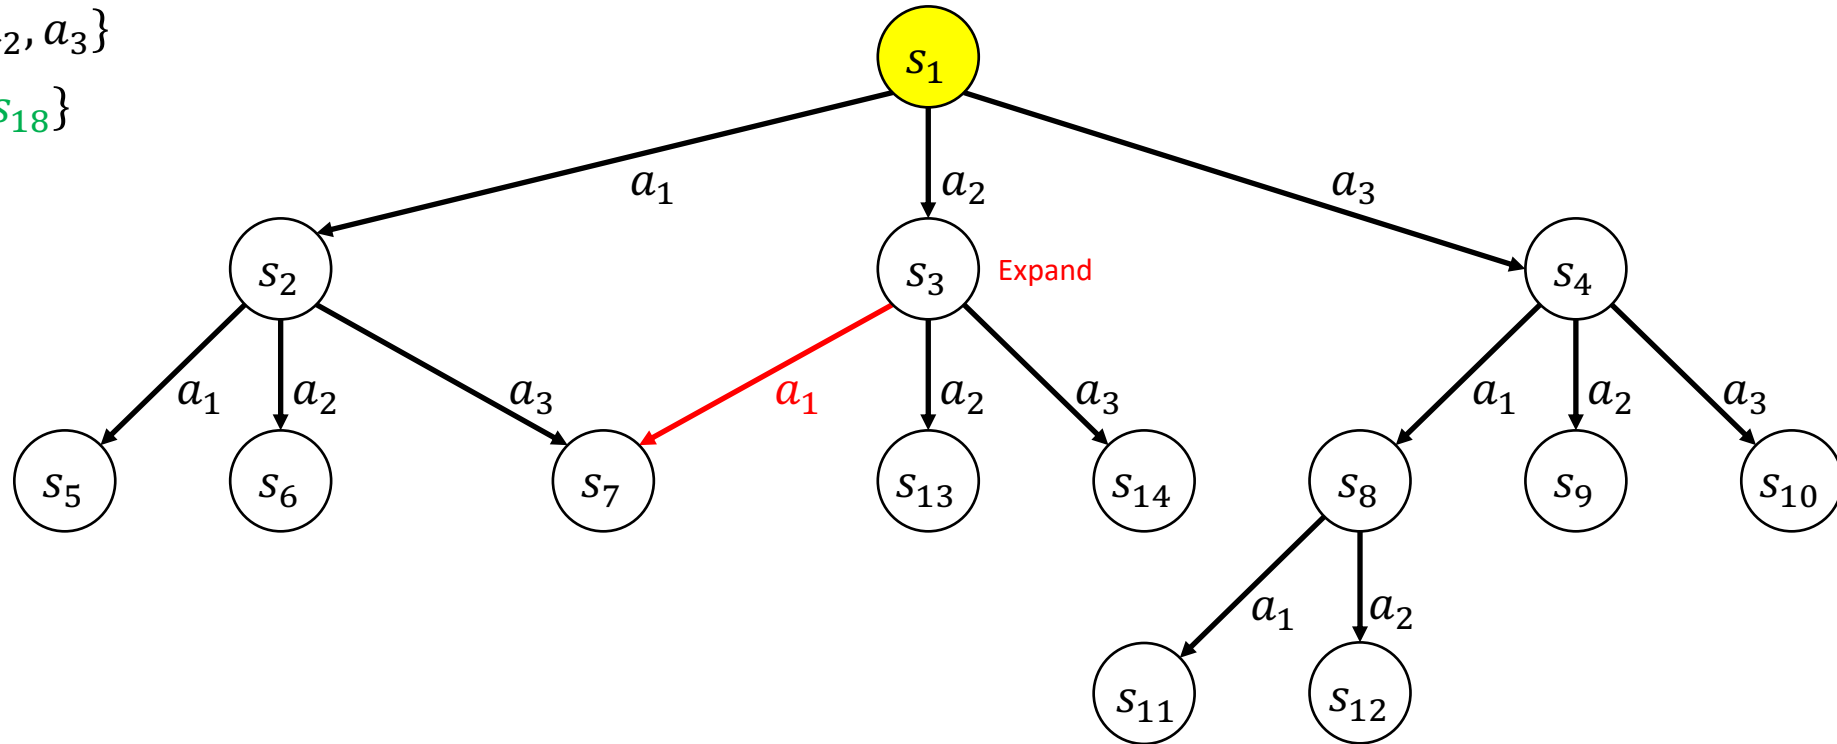
Classical Planning

Technique: A* State-Space Search (11)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



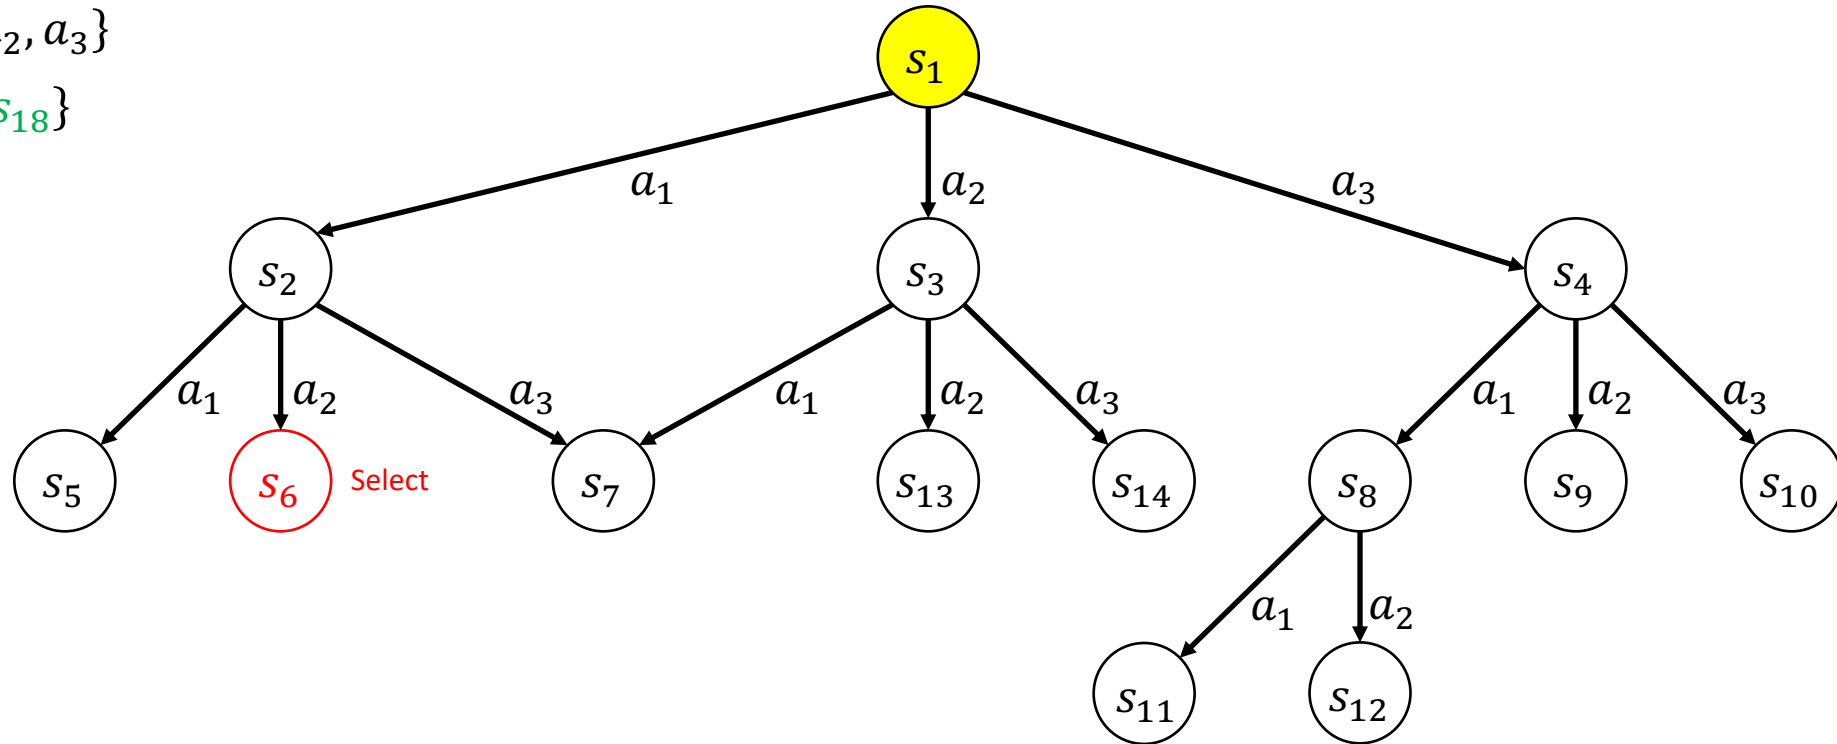
Classical Planning

Technique: A* State-Space Search (12)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



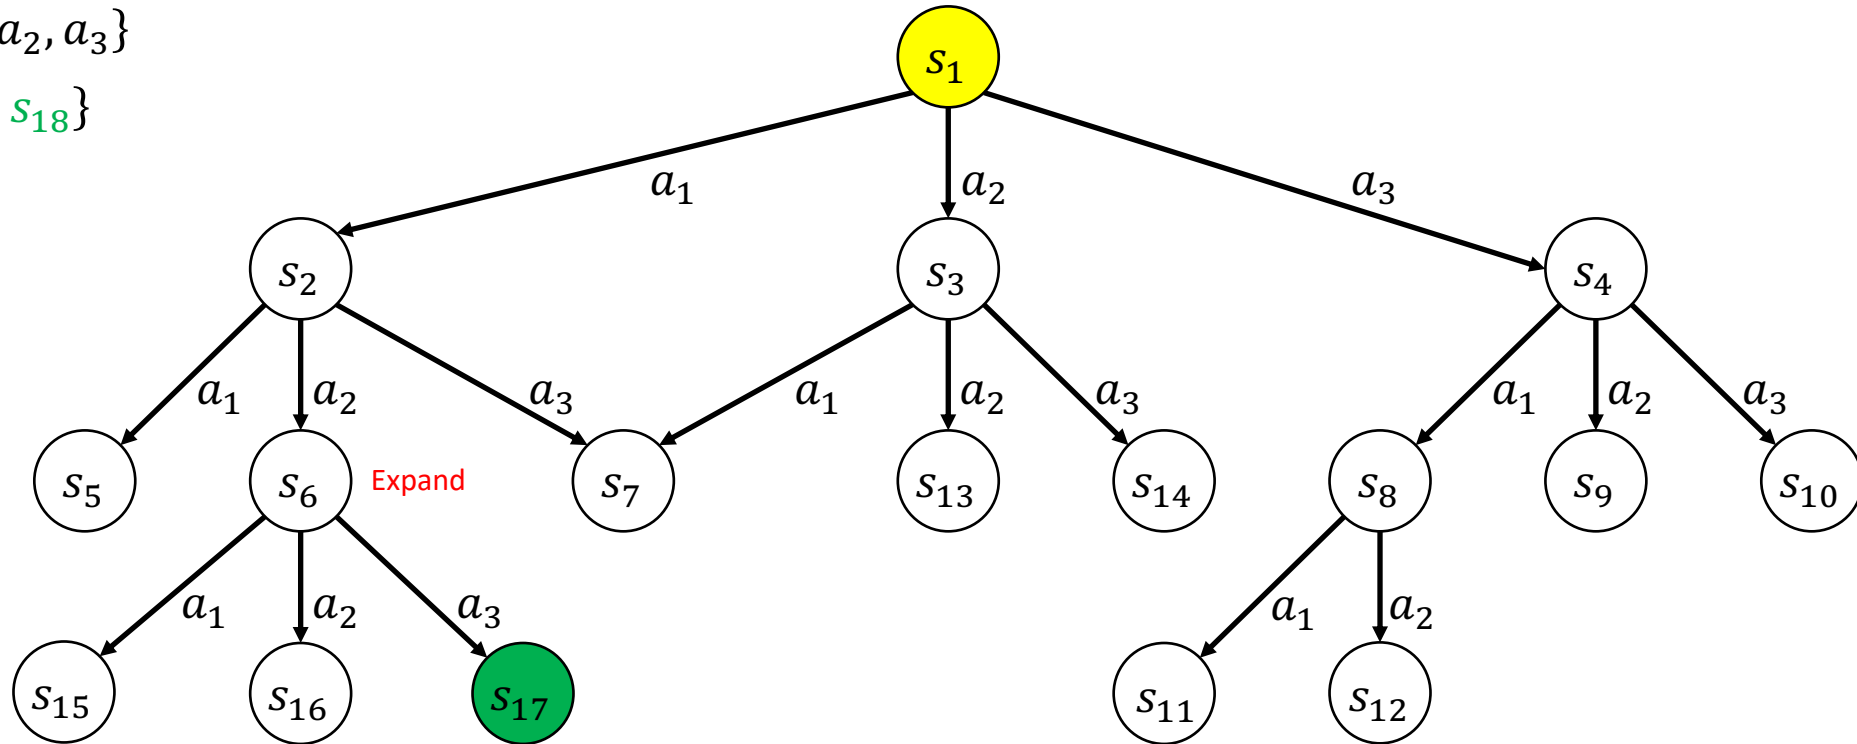
Classical Planning

Technique: A* State-Space Search (13)

$\mathcal{S} = \{s_1, s_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G = \{s_{17}, s_{18}\}$



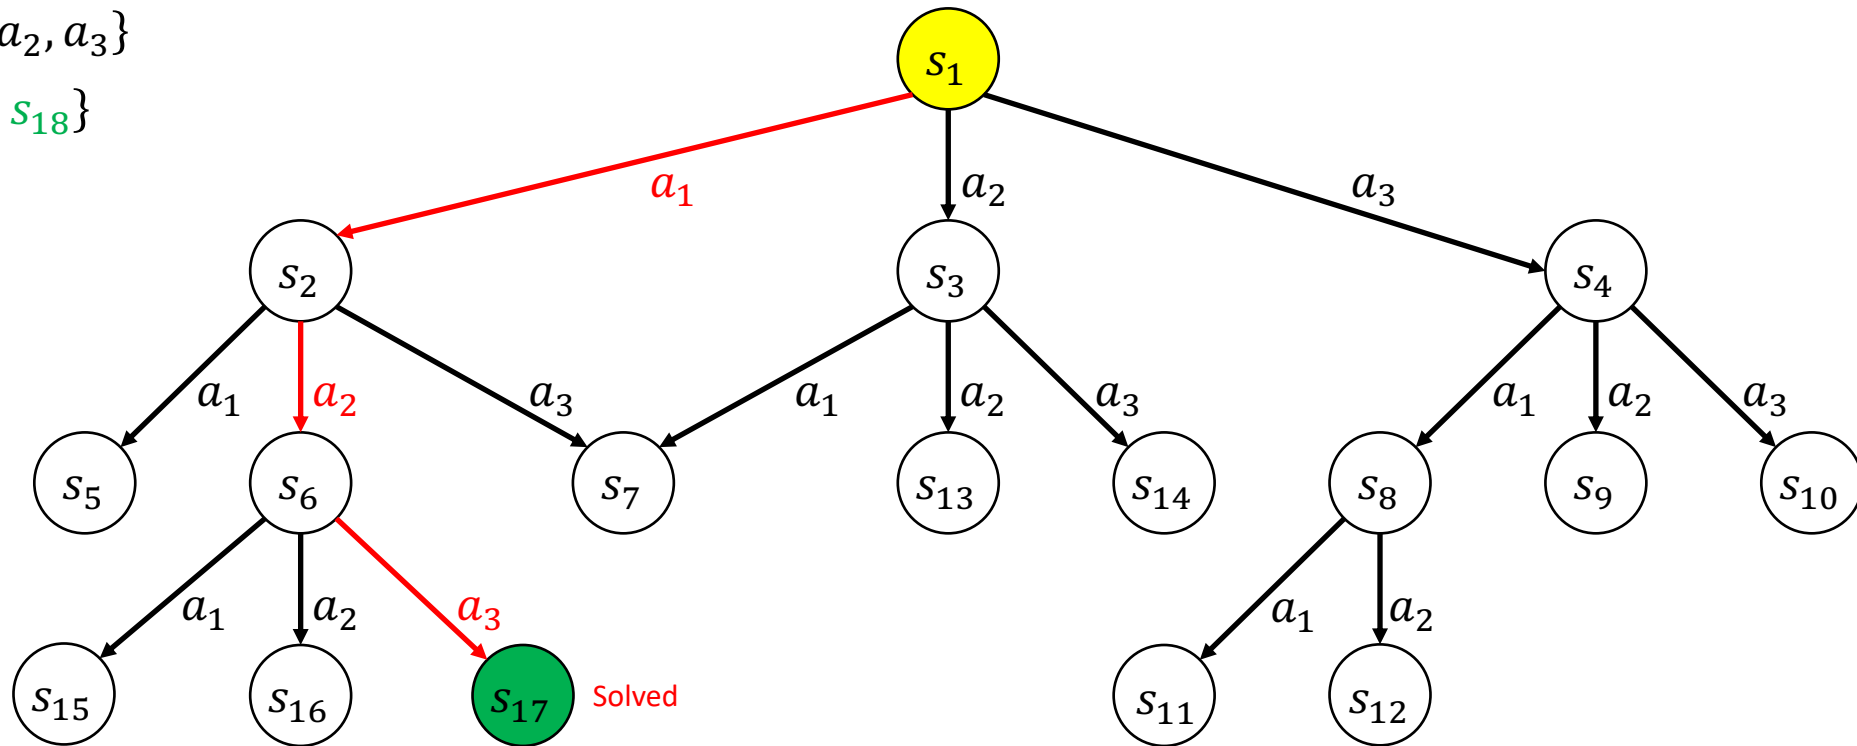
Classical Planning

Technique: A* State-Space Search (14)

$$\mathcal{S} = \{s_1, s_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2, a_3\}$$

$$G = \{s_{17}, s_{18}\}$$



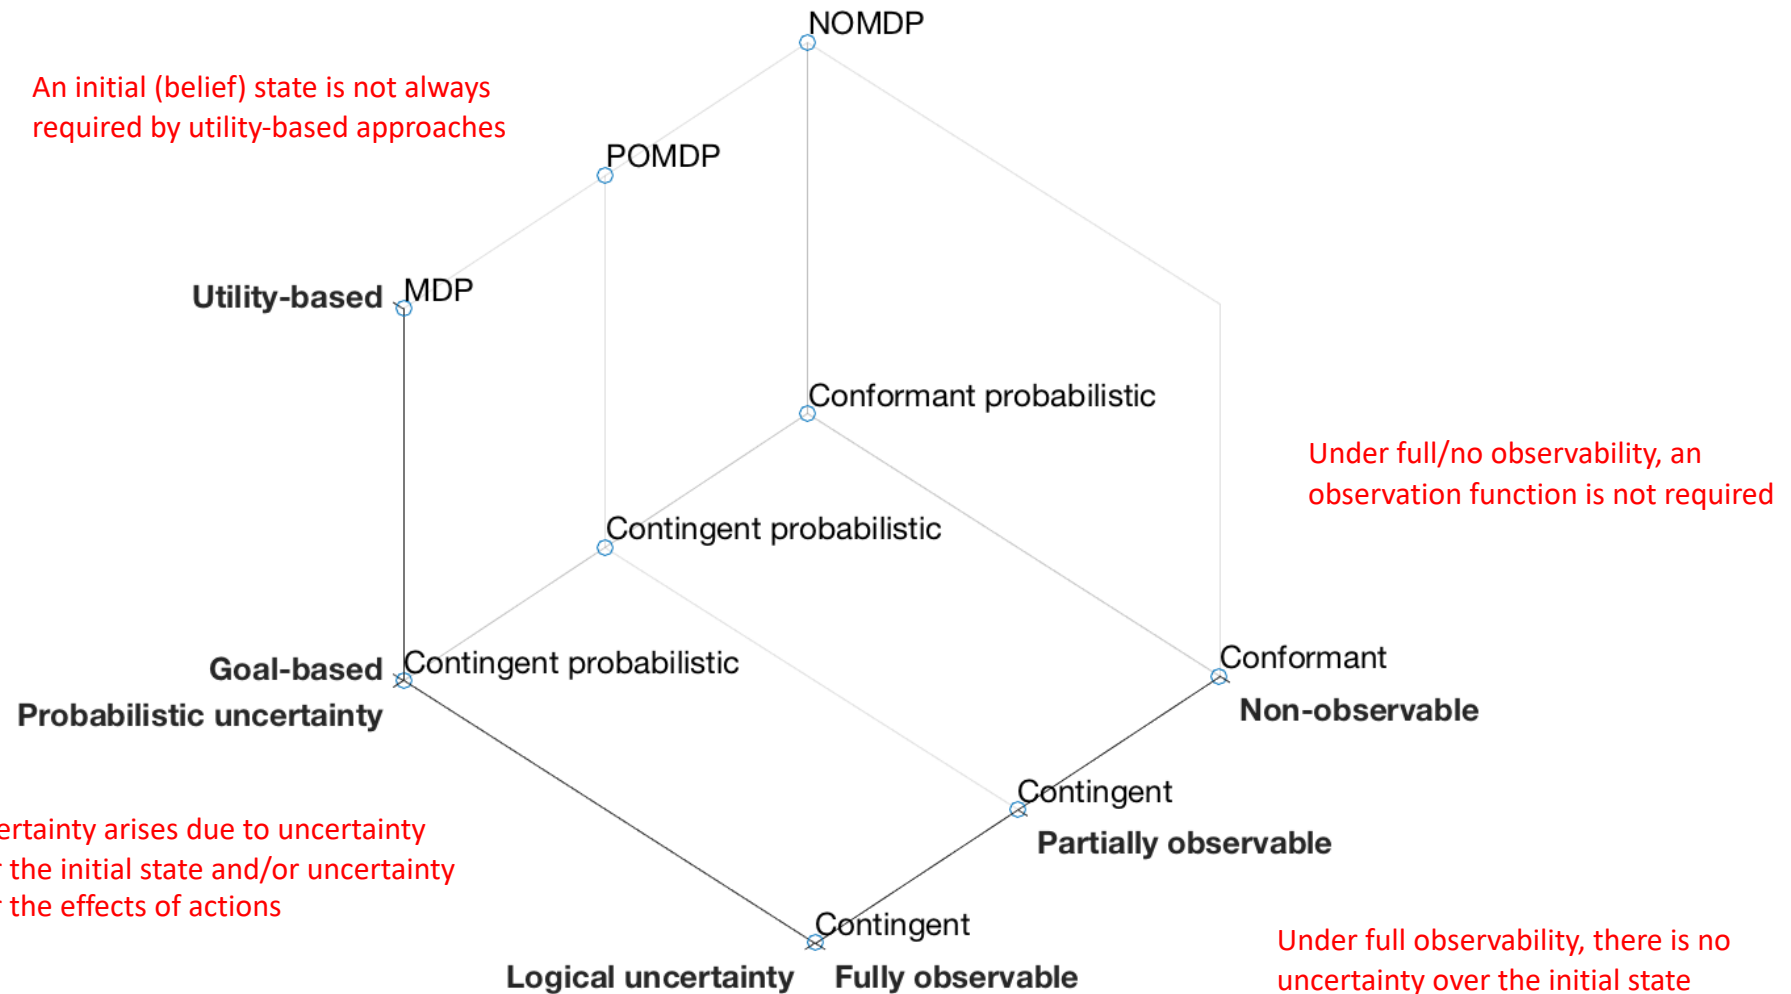
Plan $\pi = (a_1, a_2, a_3)$

Classical Planning

Pros & Cons

- - Generally not reflective of the real-world
 - In practice, plans may be sub-optimal or incomplete (e.g. if, in reality, actions are not actually deterministic)
- + Comparatively easy to describe
 - Standard PDDL action descriptions are fairly compact and intuitive
- + State-of-the-art classical planners (e.g. FF) can scale to very large/complex problems
 - Benefit of making strong (unrealistic?) assumptions about properties of the task environment
- + Richer planning problems can be “translated” into classical planning
 - Arguably, classical planning remains an active area of research largely for this reason

Planning Under Uncertainty



Conformant Planning


Definition

- Conformant planning problem $(\mathcal{S}, \mathcal{A}, b_1, T, G)$
 - Initial **belief state** $b_1 \in B$ where $B = 2^{\mathcal{S}}$
 - **Deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ or **non-deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow 2^{\mathcal{S}}$
 - Set of goal states $G \subseteq \mathcal{S}$
- A **plan** is a **finite sequence** of actions
 - Goal is **satisfied** when $b \subseteq G$
 - Must work for **any** possible initial state and/or successor state **without relying on observations**
- T is **extended** to belief states
 - $T(b, a) = \{T(s, a) \mid s \in b\}$ if T is **deterministic**
 - $T(b, a) = \{s' \in T(s, a) \mid s \in b\}$ if T is **non-deterministic**

Conformant Planning

Example (1)

$\pi = (\dots)$




Execute **Move(Right)**

Conformant Planning

Example (2)

$\pi = (\text{Move(Right)}, \dots)$




Execute **Move(Right)**

Conformant Planning

Example (3)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \dots)$




Execute **Move(Right)**

Conformant Planning

Example (4)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \dots)$




Execute **Move(Up)**

Conformant Planning

Example (5)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Up}), \dots)$




Execute **Move(Up)**

Conformant Planning

Example (6)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \dots)$




Execute **Move(Up)**

Conformant Planning

Example (7)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \dots)$




Execute **Move(Left)**

Conformant Planning

Example (8)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \text{Move}(\text{Left}), \dots)$




Execute **Move(Down)**

Conformant Planning

Example (9)

$\pi = (\text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Right}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \text{Move}(\text{Up}), \text{Move}(\text{Left}), \text{Move}(\text{Down}))$




Goal achieved

Conformant Planning

Solution Form

1. Finite sequence of actions $p = (a_1, a_2, \dots, a_n)$ from initial belief state $b_1 \in B$



3. Policy graph (B', E, b_1, π_B)
 - Directed edges $E \subseteq B' \times B'$
 - Rooted at b_1
 - Nodes labelled by π_B

2. Partial **belief-based** policy $\pi_B : B' \rightarrow \mathcal{A}$ with $B' \subseteq B$ such that π_B is closed with respect to $b_1 \in B$
 - If $b \in B$ is reachable from b_1 via π_B and T , then π_B is closed with respect b_1 if and only if $b \in B'$

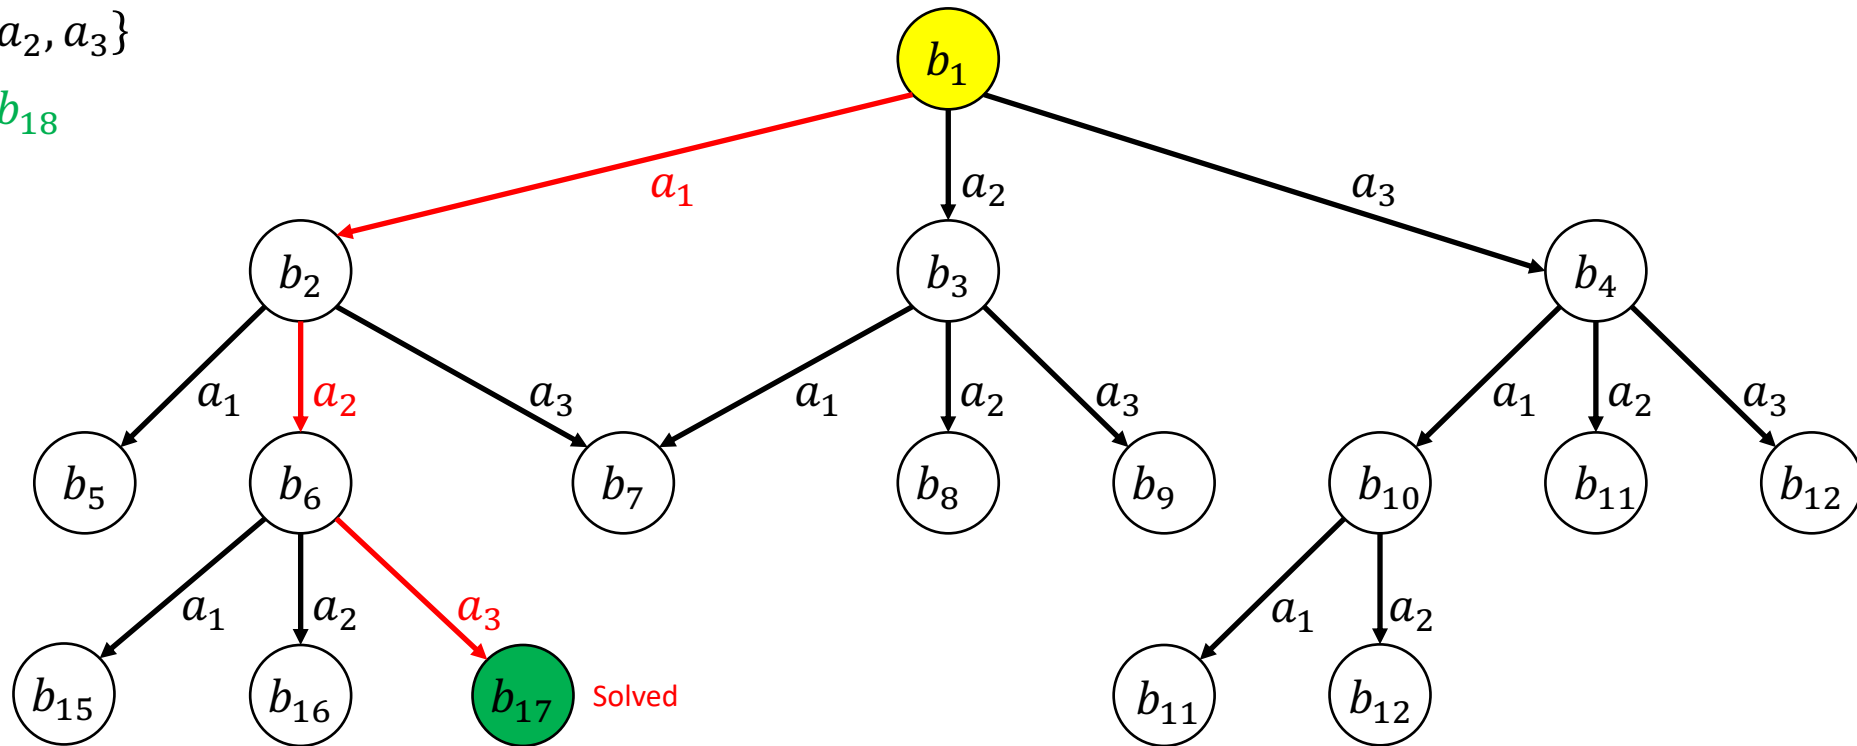
Conformant Planning

Technique: A* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2, a_3\}$

$G \ni b_{17}, b_{18}$



Plan $\pi = (a_1, a_2, a_3)$

Contingent Planning


Definition

- Contingent planning problem with **full observability** $(\mathcal{S}, \mathcal{A}, s_1, T, G)$
 - Initial **state** $s_1 \in \mathcal{S}$
 - **Non-deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow 2^{\mathcal{S}}$
 - Set of goal states $G \subseteq \mathcal{S}$
- Contingent planning problem with **partial observability** $(\mathcal{S}, \mathcal{A}, b_1, T, O, G)$
 - Initial **belief state** $b_1 \in B$ where $B = 2^{\mathcal{S}}$
 - **Deterministic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$ or non-deterministic transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow 2^{\mathcal{S}}$
 - **Deterministic** observation function $O : \mathcal{S} \rightarrow \mathcal{P}$ or **non-deterministic observation** function $O : \mathcal{S} \rightarrow 2^{\mathcal{P}}$
 - Set of goal states $G \subseteq \mathcal{S}$
- A **plan** is a **branching** structure (e.g. a **tree** or **graph**)
 - Must work for **any** possible percept
 - May contain **cycles**, but sometimes an **acyclic requirement** is imposed

Contingent Planning

Solution Form

2. Partial belief-based policy $\pi_B : B' \rightarrow \mathcal{A}$ with $B' \subseteq B$ such that π_B is closed with respect to $b_1 \in B$



3. Policy graph $(B', \mathcal{P}, E, b_1, \pi_B)$

- Directed **multi-edges** $E \subseteq B' \times \mathcal{P} \times B'$
- Rooted at b_1
- Nodes labelled by π_B


Contingent Planning

Technique: AO* Belief-Space Search

$$B = \{b_1, b_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

$$G \ni b_{16}, b_{17}, b_{18}$$




Contingent Planning

Technique: AO* Belief-Space Search

$$B = \{b_1, b_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

$$G \ni b_{16}, b_{17}, b_{18}$$




Contingent Planning

Technique: AO* Belief-Space Search

$$B = \{b_1, b_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

$$G \ni b_{16}, b_{17}, b_{18}$$



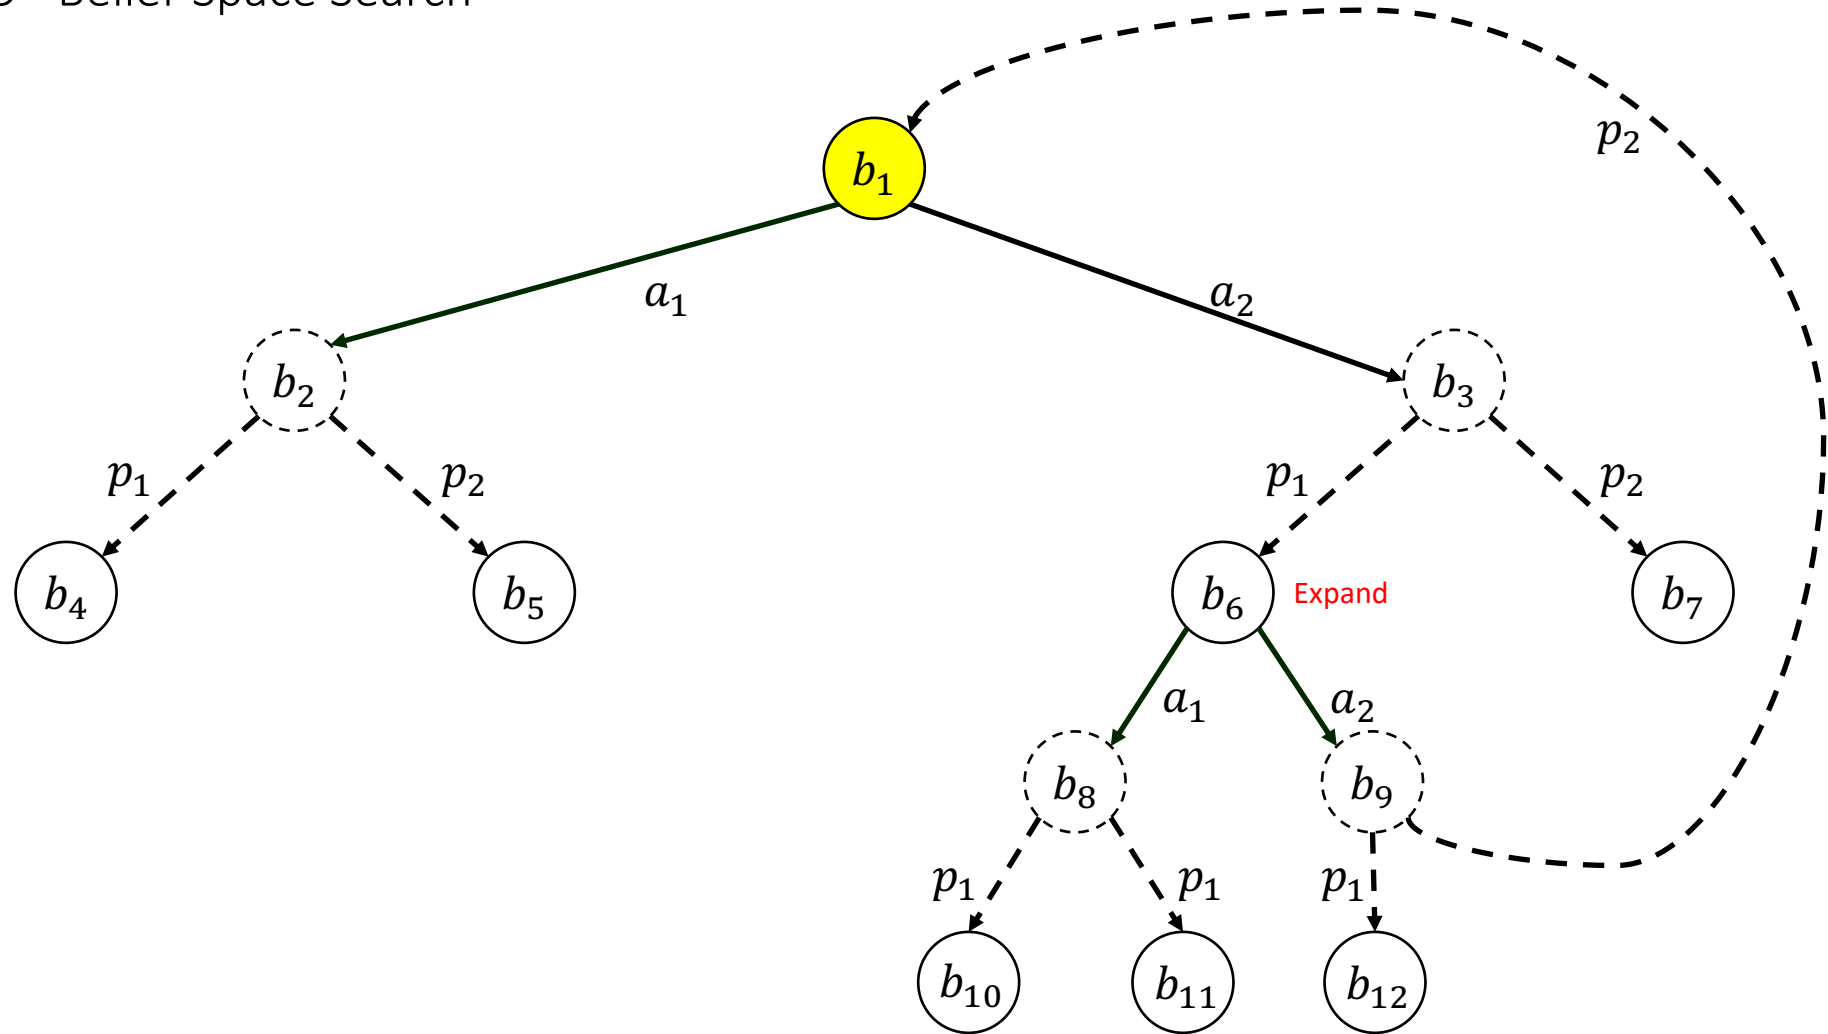
Contingent Planning

Technique: AO* Belief-Space Search

$$B = \{b_1, b_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

$$G \ni b_{16}, b_{17}, b_{18}$$



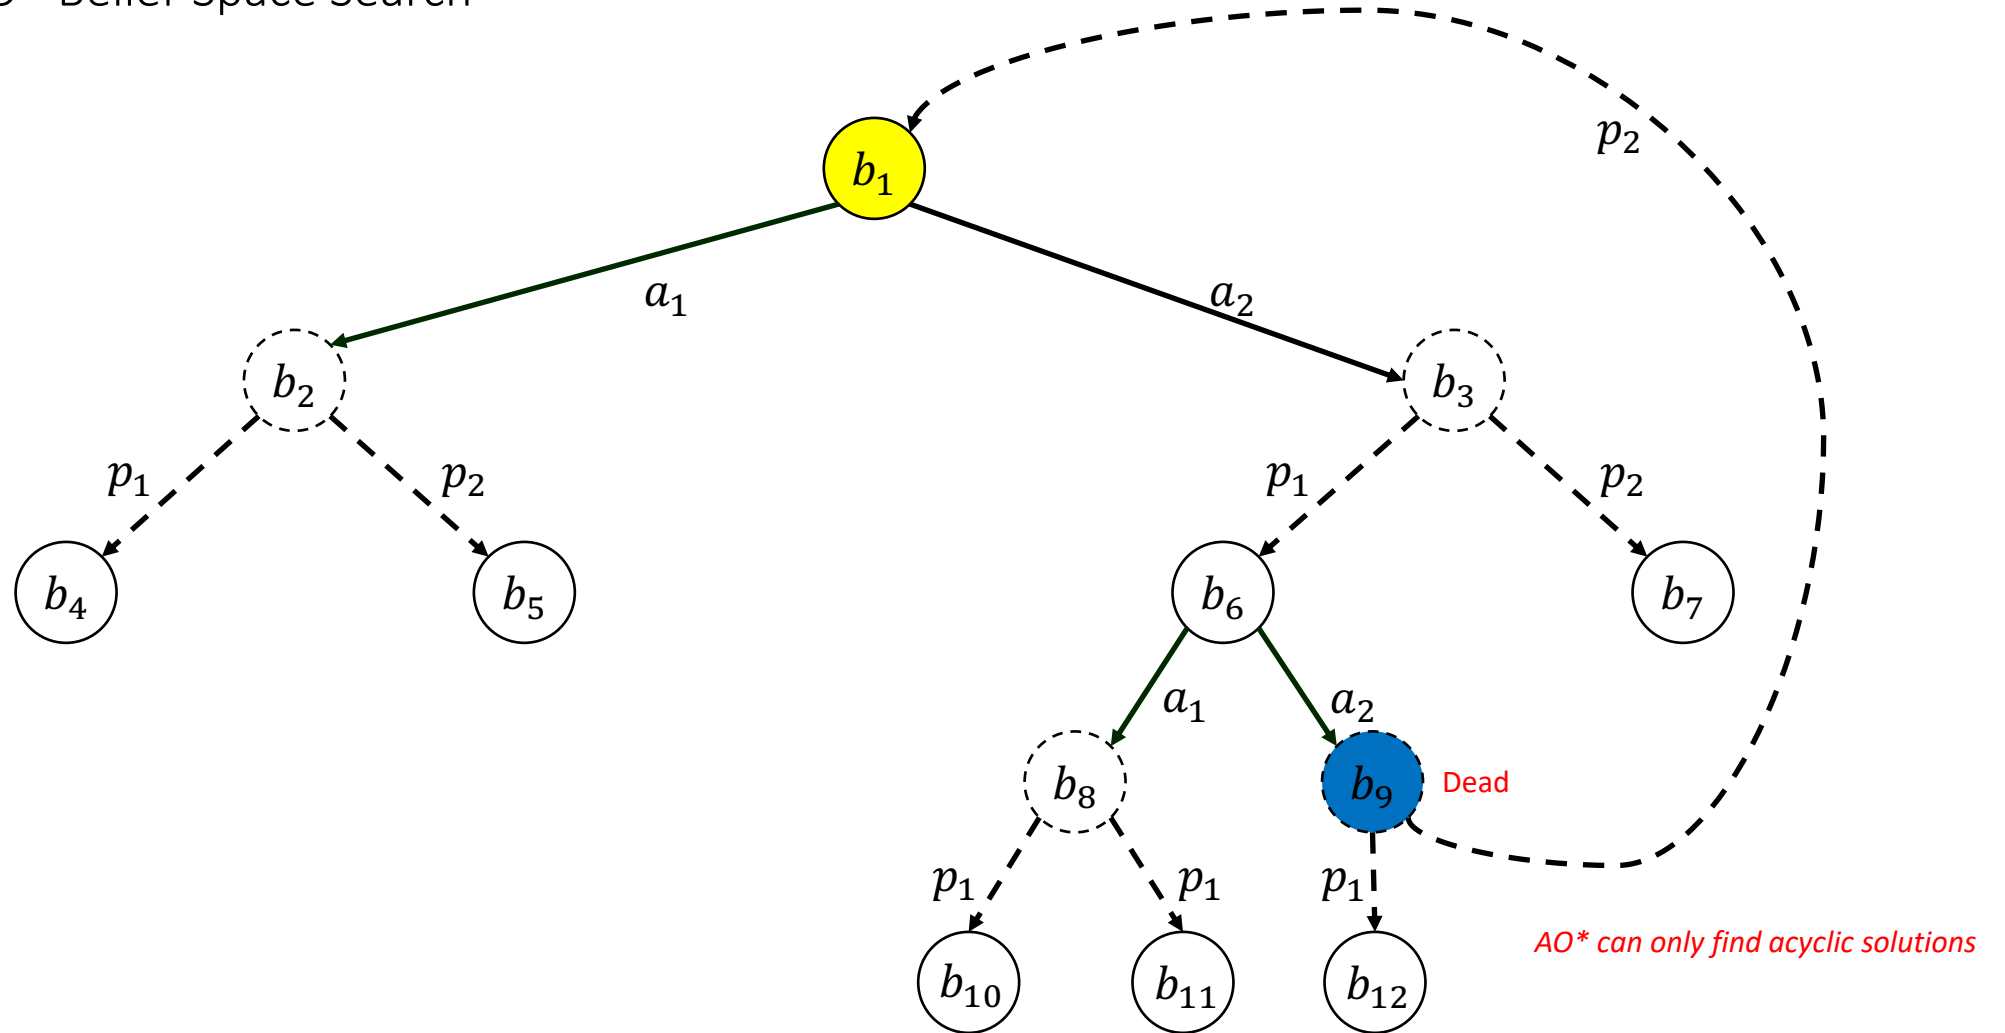
Contingent Planning

Technique: AO* Belief-Space Search

$$B = \{b_1, b_2, \dots\}$$

$$\mathcal{A} = \{a_1, a_2\}$$

$$G \ni b_{16}, b_{17}, b_{18}$$



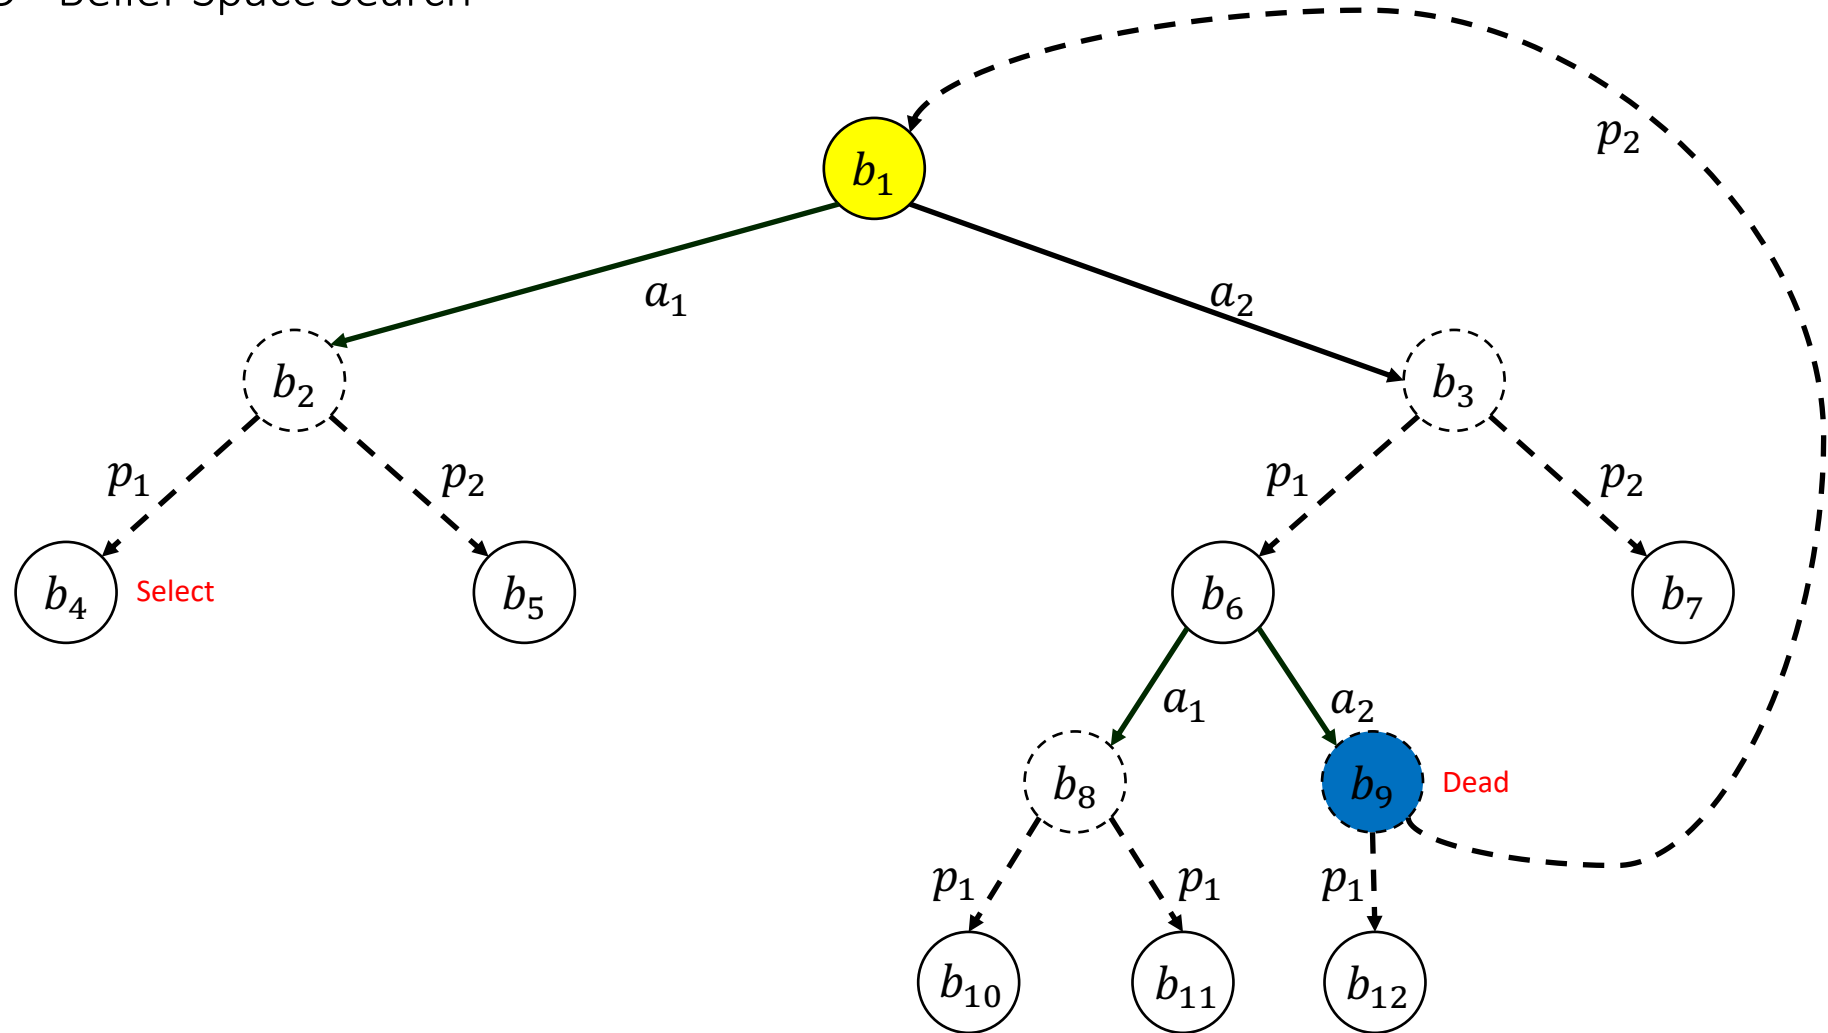
Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$




Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$



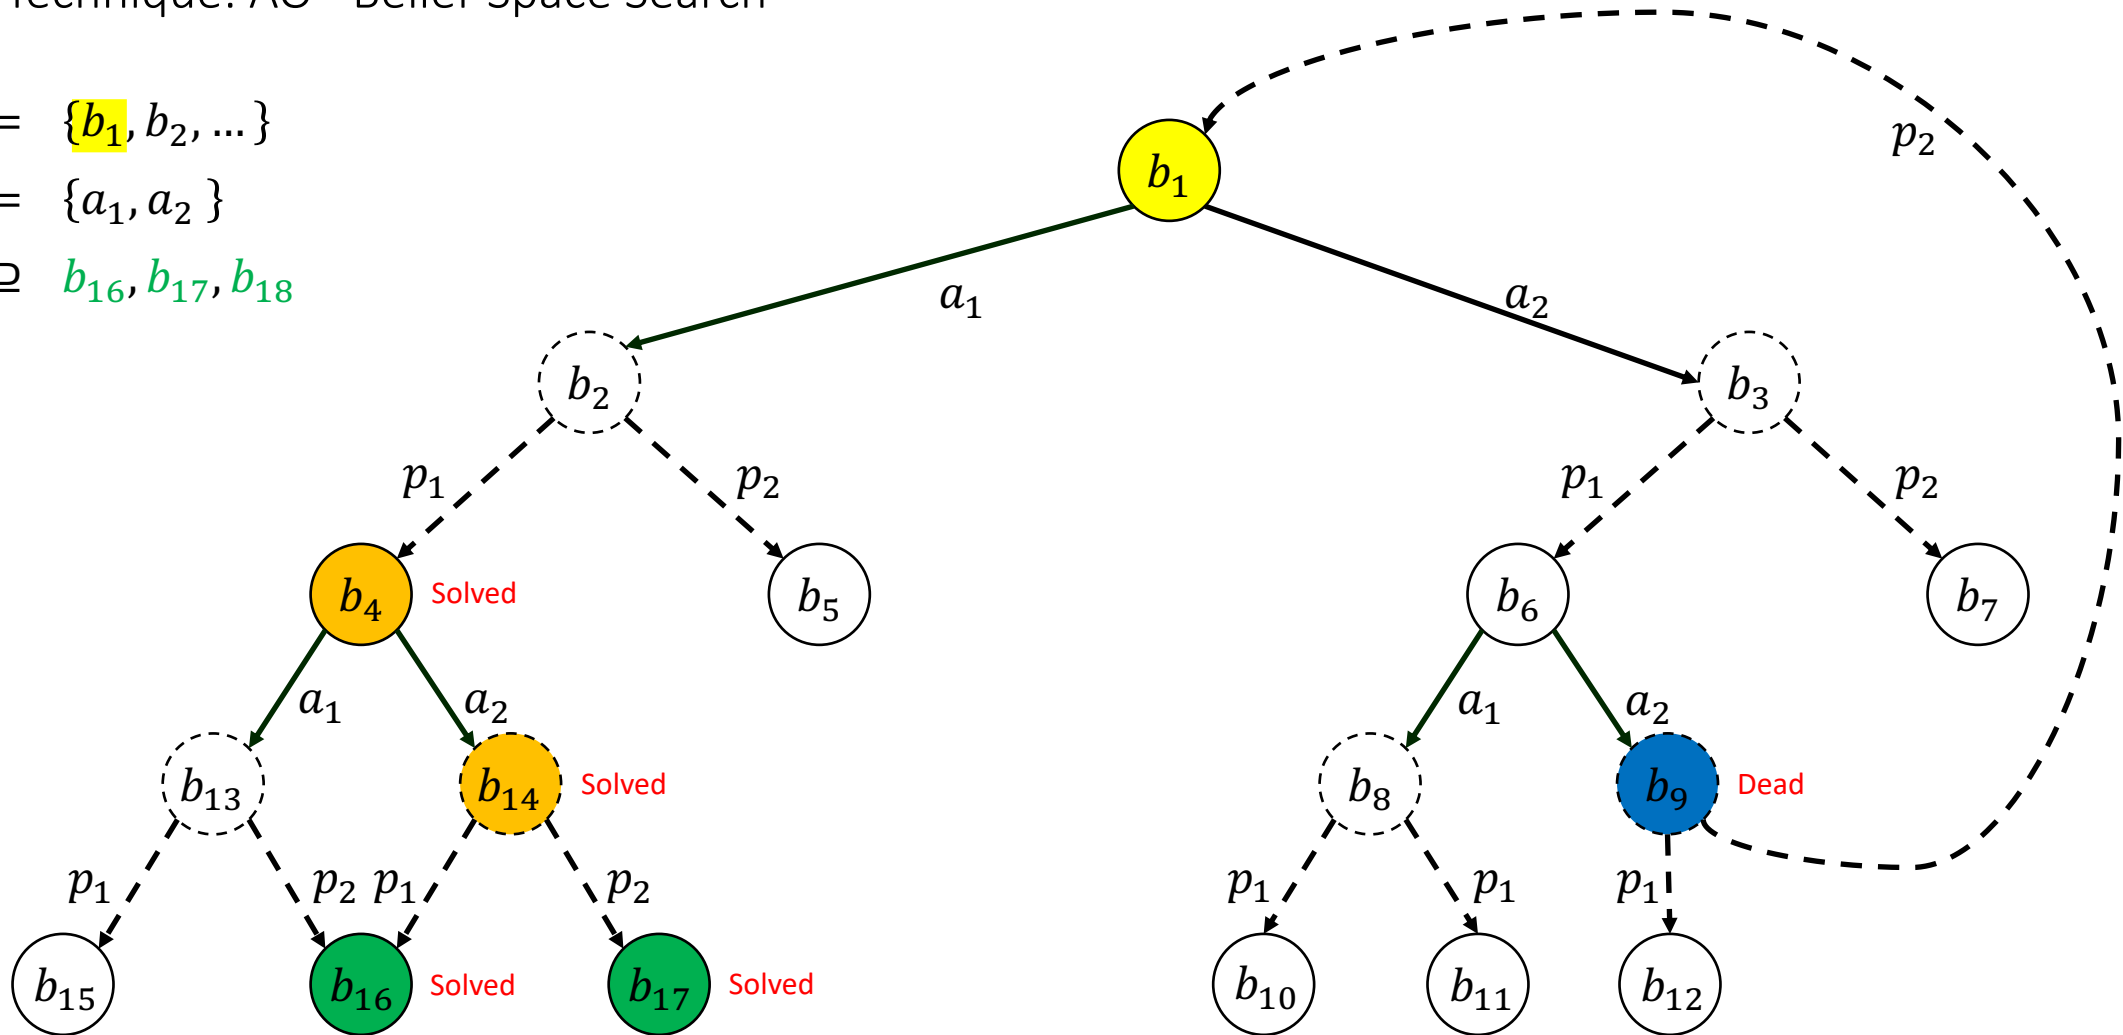
Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$




Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$




Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$



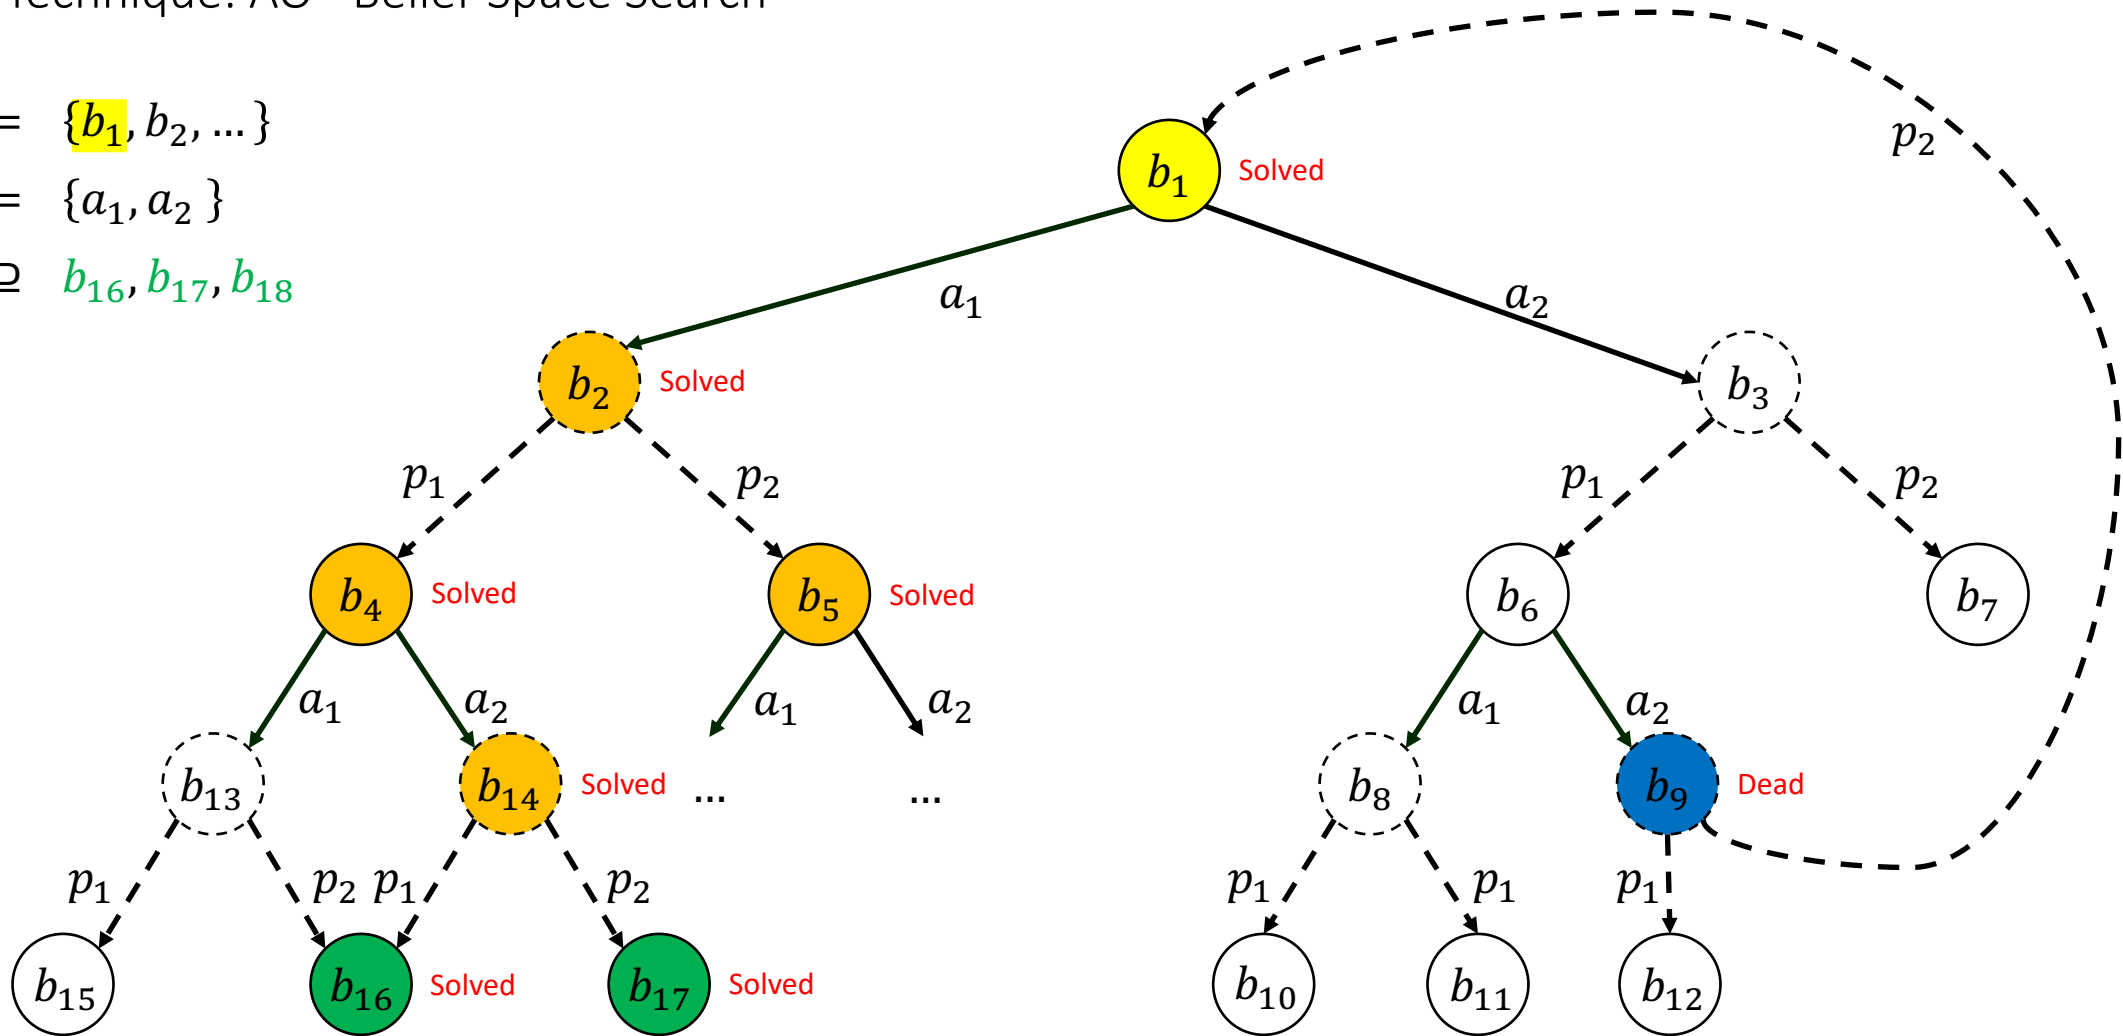
Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$




Contingent Planning

Technique: AO* Belief-Space Search

$B = \{b_1, b_2, \dots\}$

$\mathcal{A} = \{a_1, a_2\}$

$G \ni b_{16}, b_{17}, b_{18}$



Markov Decision Process (MDP)

Definition (Fully Observable)

- **Finite-horizon** MDP $(\mathcal{S}, \mathcal{A}, T, R, h)$

- **Stochastic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$
- **Utility** function $R : \mathcal{S} \rightarrow \mathbb{R}$
- **Horizon** $h \in \mathbb{N}$

- Infinite-horizon, **discounted-reward** MDP $(\mathcal{S}, \mathcal{A}, T, R, \gamma)$

- **Stochastic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$
- **Utility** function $R : \mathcal{S} \rightarrow \mathbb{R}$
- **Discount factor** $0 \leq \gamma < 1$

- Indefinite-horizon, **goal-oriented** MDP $(\mathcal{S}, \mathcal{A}, s_1, T, G)$

- Initial **state** $s_1 \in \mathcal{S}$
- **Stochastic** transition function $T : \mathcal{S} \times \mathcal{A} \rightarrow \Pi(\mathcal{S})$
- Set of **goal states** $G \subseteq \mathcal{S}$

Equivalent definitions for partially observable MDPs (POMDPs) and non-observable MDPs (NOMDPs)

Type of MDP	Type of planning
Goal-oriented MDP	Contingent probabilistic planning with full observability
Goal-oriented POMDP	Contingent probabilistic planning with partial observability
Goal-oriented NOMDP	Conformant probabilistic planning

Online Planning

- What we have discussed so far is known as **offline planning**
 - Generate a **complete plan** and then execute in full
- An alternative is **online planning**
 - Generate an **incomplete policy**, execute until it is necessary to plan again, and repeat
 - Special case: generate “**next action**”, execute, and repeat
 - An **online planner** is essentially an **agent**, since it must run in an **environment** and **execute actions**
- **Incomplete policy**
 - A partial state- or belief-based policy π is **incomplete** if π is **not closed** with respect to initial (belief) state $b_1 \in B$
- Plan-act **procedure**
 1. Execute $\pi_B : B' \rightarrow \mathcal{A}$ until some belief state $b \in B \setminus B'$ is encountered
 2. Update π_B to π_B' via planning such that $\pi_B' : B'' \rightarrow \mathcal{A}$ with $B'' \supset B'$
 3. **Return to 1.**

Online Planning

Technique: Classical Replanning

Algorithm: Classical replanning

Input: Goal-oriented MDP $(\mathcal{S}, \mathcal{A}, s_1, T, G)$

```
1  $s \leftarrow s_1$ 
2 for each  $s' \in \mathcal{S}$  do
3   for each  $a \in \mathcal{A}$  do
4      $T'(s', a) \leftarrow \operatorname{argmax}_{s'' \in \mathcal{S}} T(s', a, s'')$  // most-probable outcome
5 while  $s \notin G$  do
6   if  $\pi(s) = \text{undefined}$  then
7      $\pi' \leftarrow \text{CLASSICAL-PLAN}(\mathcal{S}, \mathcal{A}, s, T', G)$  // replan
8      $\pi \leftarrow \text{UPDATE}(\pi, \pi')$ 
9    $\text{EXECUTE}(\pi(s))$ 
10   $s \leftarrow \text{SENSE}()$ 
```

Full observability

Many other determinization methods
(e.g. all outcomes with costs as
inverse transition probabilities)

Automated Planning & Agents

Some Reflections

- A **plan/policy** can be thought of as an **agent function**
 - See panel discussion on **BDI vs. MDPs** at MSDM'15 workshop @ AAMAS'15: “encode **everything** in an MDP”
 - A plan/planner provides a very **limited** view of agents (e.g. no multi-tasking, no exogenous events, no complex reasoning)
- Agent programs are **compatible** with automated planning
 - Plans can be **generated** by an automated planner and then **executed** by an agent
- Potential for plan **misuse**
 - For example, **interleaving** the execution of multiple classical plans may fail due to **conflicting action-effects**
 - Representing a plan as a **policy** helps to avoid this issue

Automated Planning & Agents

Other Approaches

- Partial-order (or least-commitment) planning
 - Plan is a **partially ordered set** of actions – no convenient state- or belief-based policy representation(?)
 - Difficult to extend to richer planning problems – especially those with **probabilistic uncertainty**
- Hierarchical task networks (HTN)
 - Input is a set **compound tasks** rather than a set of actions
 - Aim is to find a valid **decomposition** of those tasks into (executable) actions
 - Many successful **industrial applications**
- BDI-style plan selection
 - **Instantiation** of lifted plans and decomposition of **sub-goals**
 - Similar to HTN

This Lecture...

1. Classical planning
 - Planning languages
2. Planning under uncertainty
 - Conformant planning
 - Contingent planning
 - Markov decision processes
3. Online planning
 - Classical replanning

Next Lecture...

1. Agent-oriented programming
 - PRS
 - AgentSpeak
2. AgentSpeak
 - Logic programming
 - Syntax
 - Semantics
 - Interpreter
 - Programming

Bibliography

- Michael Wooldridge. *An Introduction to Multiagent Systems*. John Wiley & Sons, 2nd edition, 2009.

Suggested reading (Chapter 3, 4, 10, 11, and 17)

- Stuart J. Russell & Peter Norvig. [*Artificial Intelligence: A Modern Approach*](#). Prentice Hall, 3rd edition, 2009.

Suggested reading

- Hector Geffner & Blai Bonet. *A Concise Introduction to Models and Methods for Automated Planning*. Morgan & Claypool, 2013.

- Anand S. Rao. [*AgentSpeak\(L\): BDI agents speak out in a logical computable language*](#). In *Proceedings of the 7th European Workshop on Modelling Autonomous Agents in a Multi-Agent World (MAAMAW'96)*, pages 42-55, 1996.
- Rafael H. Bordini, Jomi Fred Hübner, & Michael Wooldridge. [*Programming Multi-Agent Systems in AgentSpeak Using Jason*](#). John Wiley & Sons, 2007.